



deeplearning.ai

# Case Studies

---

Why look at  
case studies?

# Outline

Classic networks:

- LeNet-5 ←
- AlexNet ←
- VGG ←

ResNet (152)

Inception



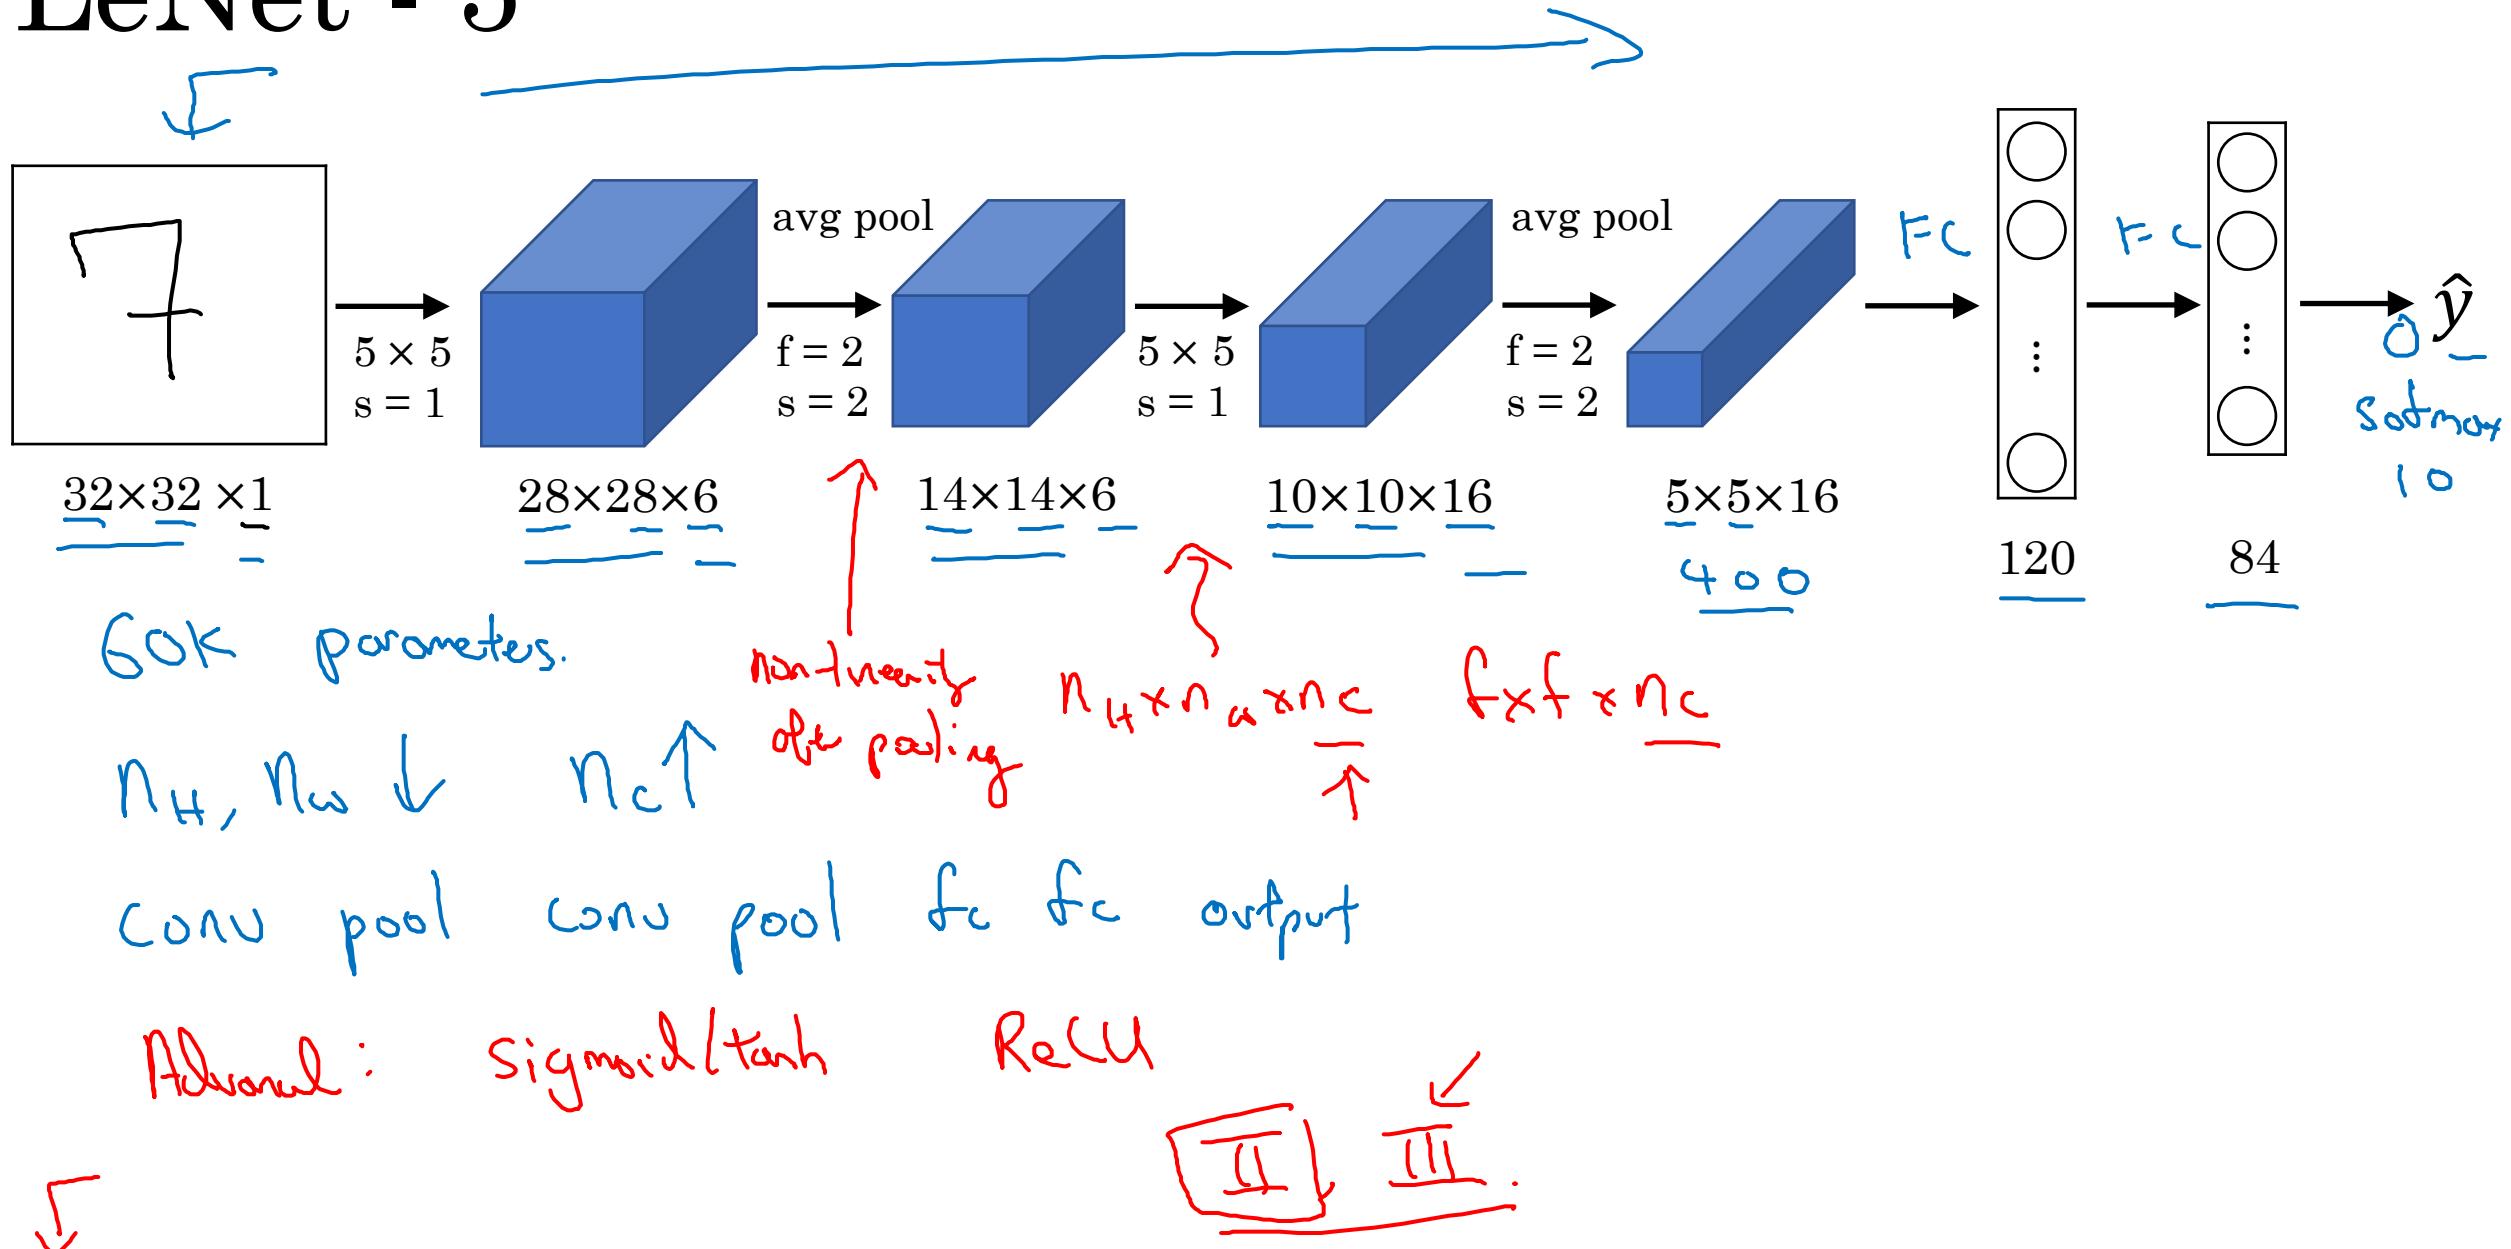
deeplearning.ai

# Case Studies

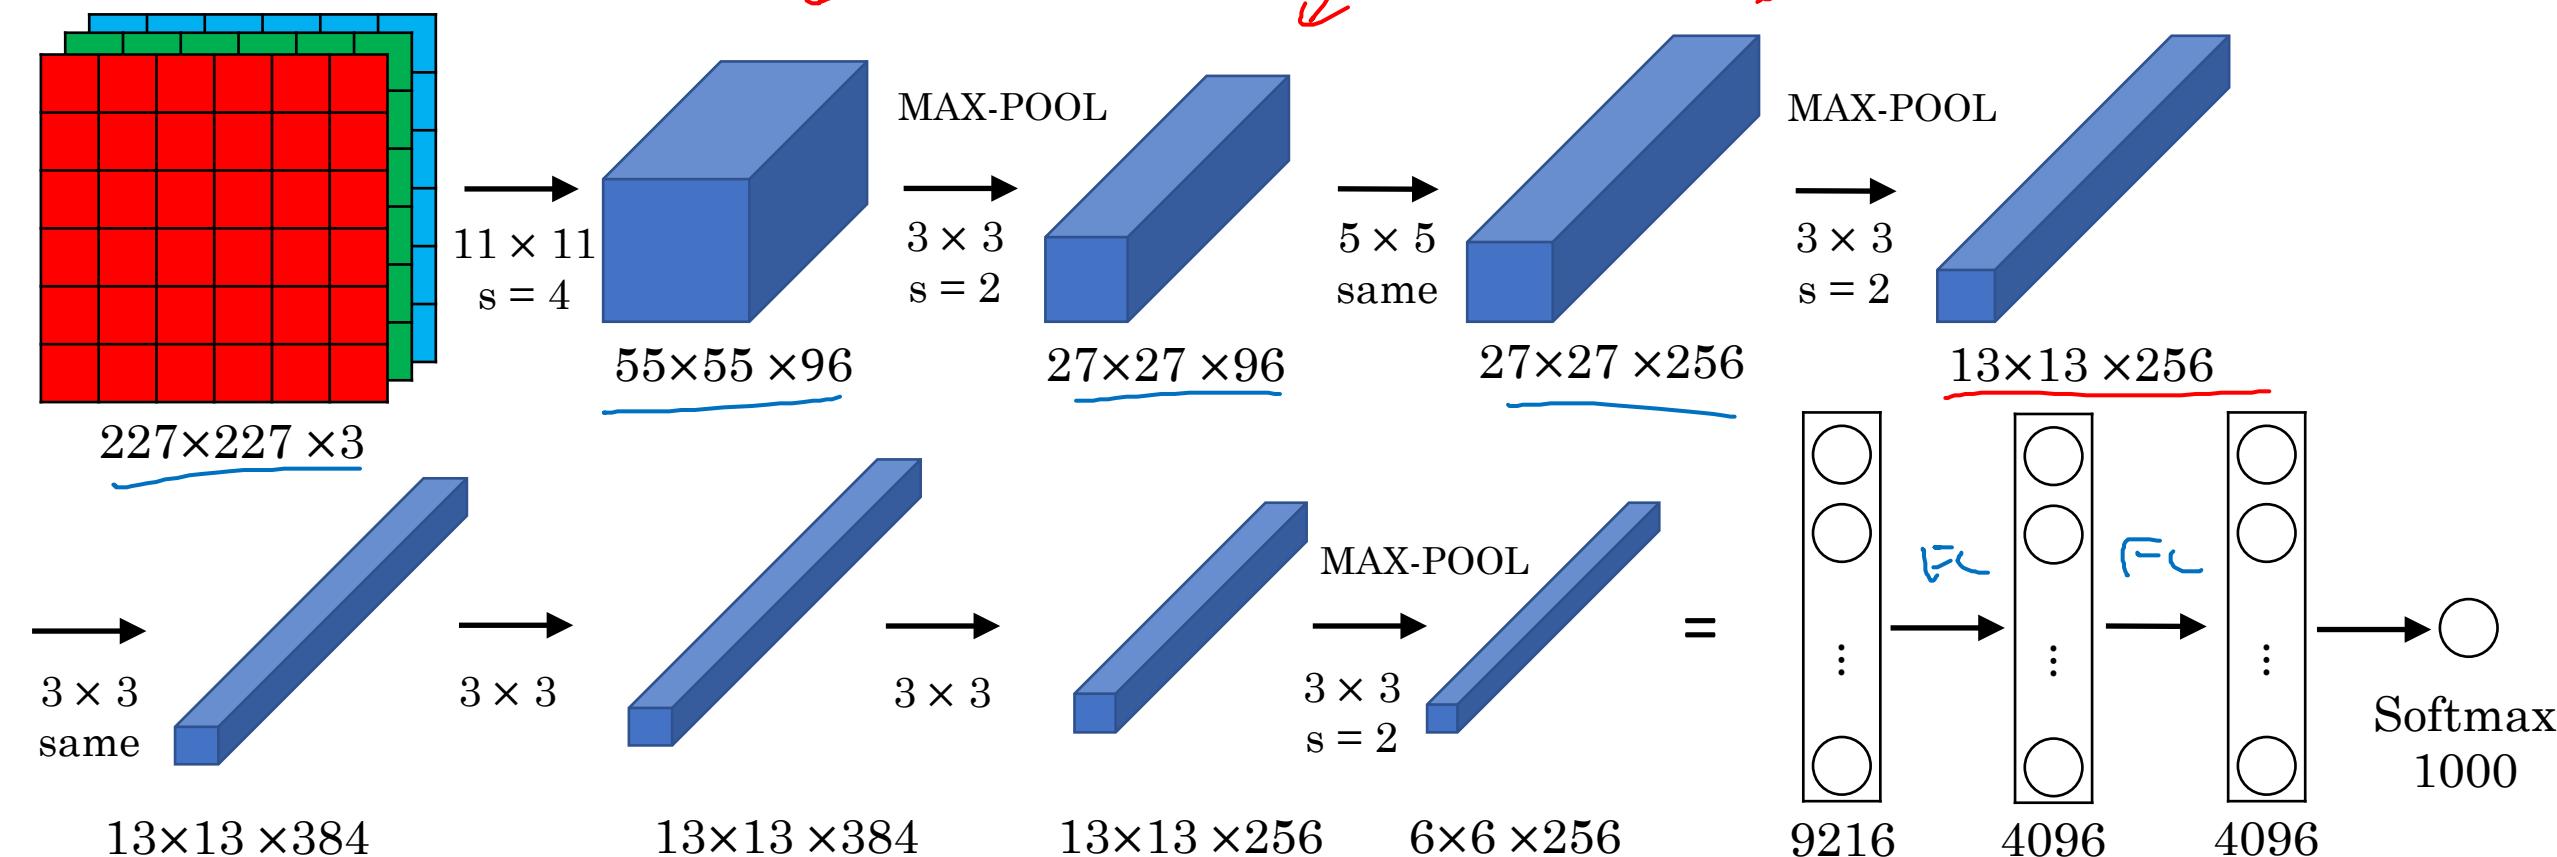
---

## Classic networks

# LeNet - 5



# AlexNet

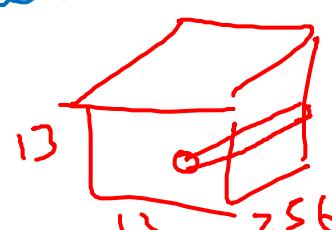


- Similar to LeNet, but much bigger.

- ReLU

- Multiple GPUs.

- Local Response Normalization (LRN)

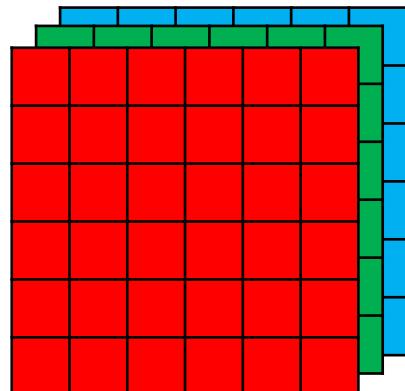


~60M Parameters

Andrew Ng

# VGG - 16

CONV =  $3 \times 3$  filter,  $s = 1$ , same



$224 \times 224$

[CONV 64]  
 $\times 2$

$224 \times 224 \times 64$  → POOL →  $112 \times 112 \times 64$

$112 \times 112 \times 64$

[CONV 128]  
 $\times 2$

$112 \times 112 \times 128$  → POOL →  $56 \times 56 \times 128$

$224 \times 224$

$56 \times 56 \times 256$  → POOL →  $28 \times 28 \times 256$  → [CONV 512]  
 $\times 3$  →  $28 \times 28 \times 512$  → POOL →  $14 \times 14 \times 512$

$14 \times 14 \times 512$  → POOL →  $7 \times 7 \times 512$  → FC 4096 → FC 4096 → Softmax 1000  
[CONV 512]  
 $\times 3$

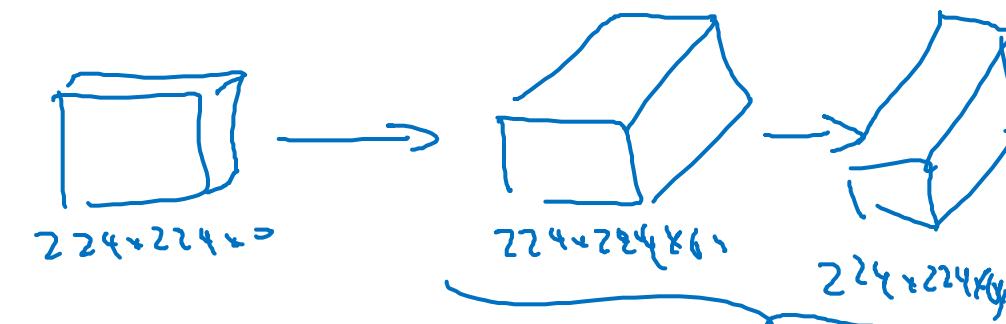
$n_h, n_w \downarrow$

$n_c \uparrow$

$\sim 38M$

# VGG-19

MAX-POOL =  $2 \times 2$ ,  $s = 2$



$112 \times 112 \times 128$  → POOL →  $56 \times 56 \times 128$

POOL



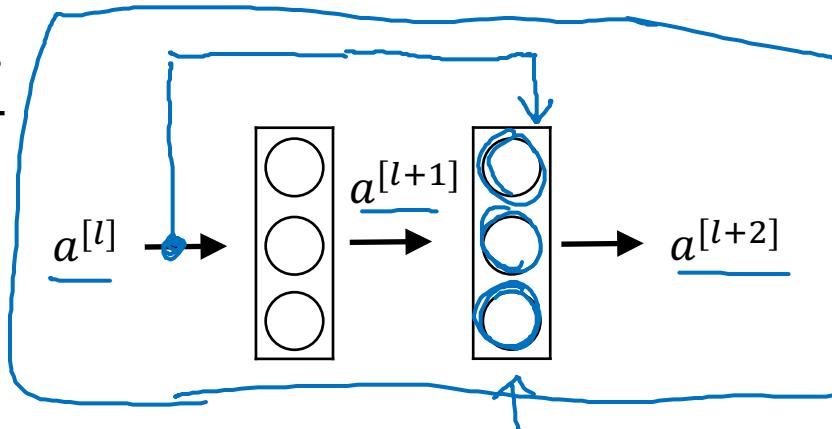
deeplearning.ai

## Case Studies

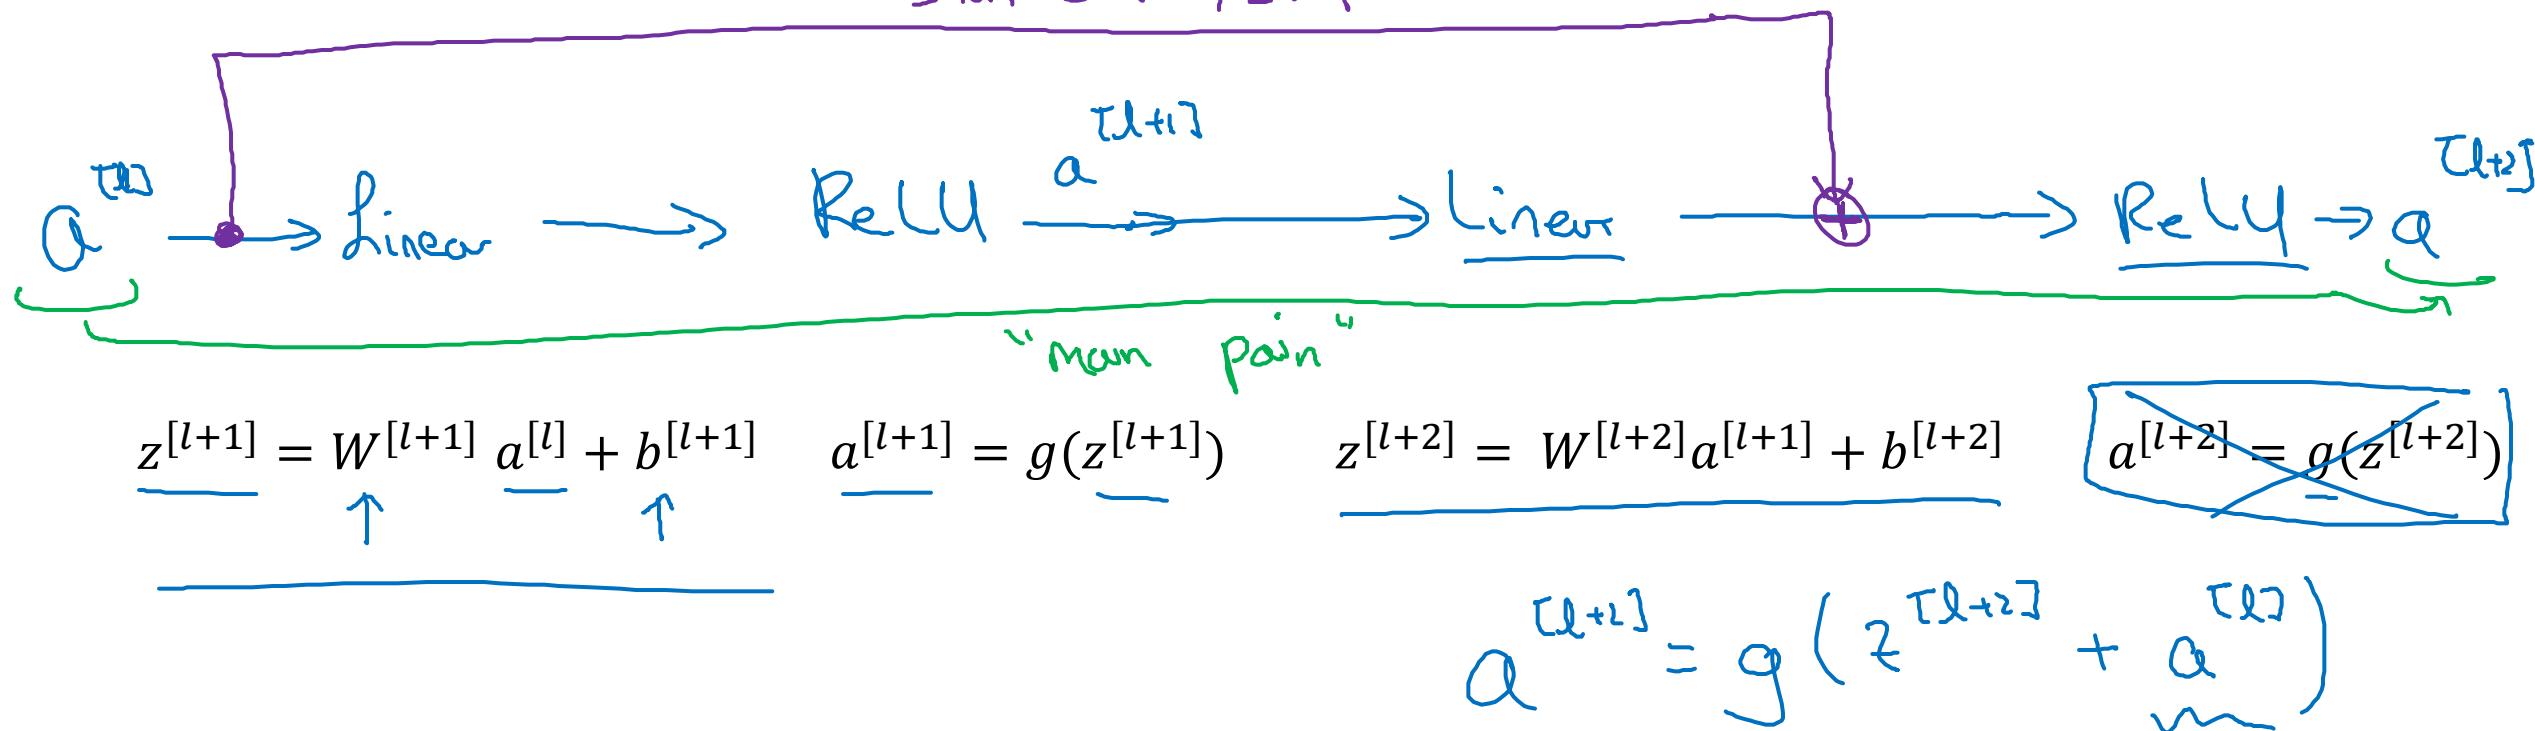
---

# Residual Networks (ResNets)

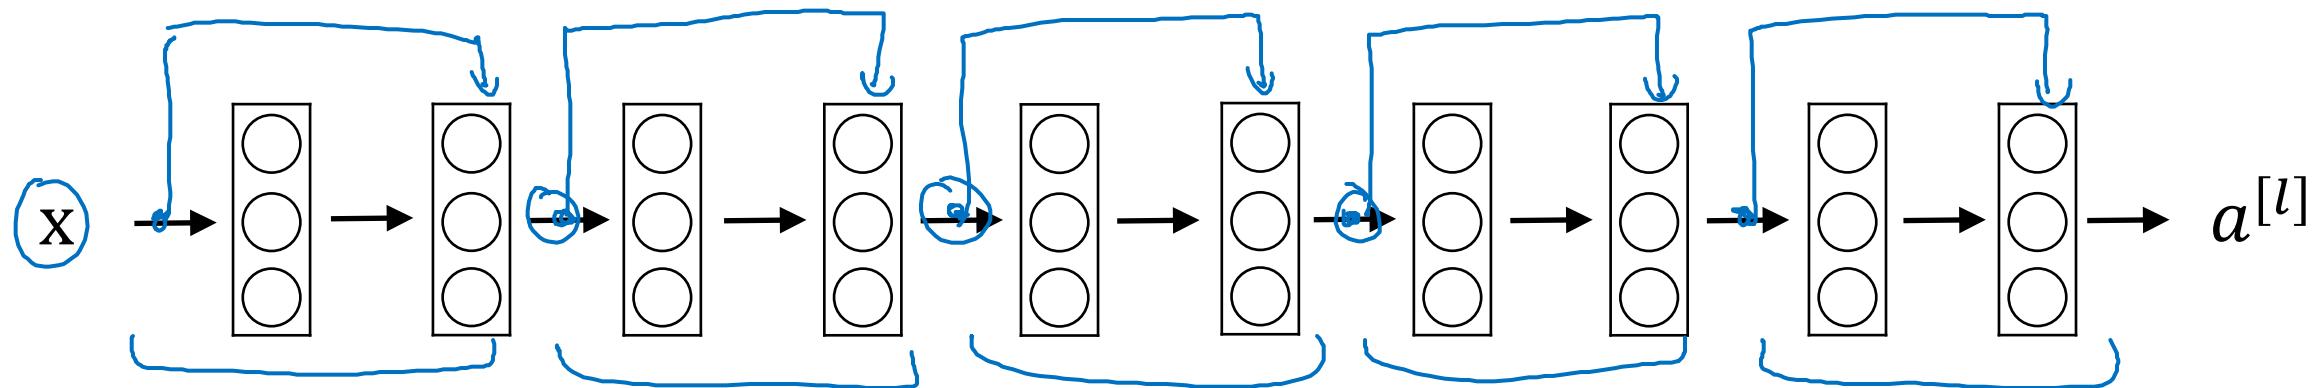
# Residual block



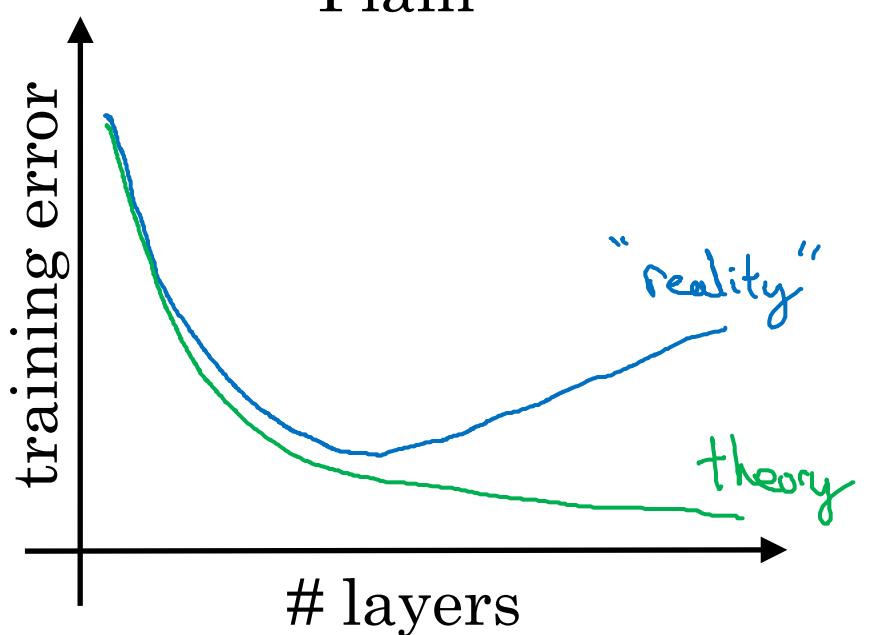
"Short cut" /skip connection



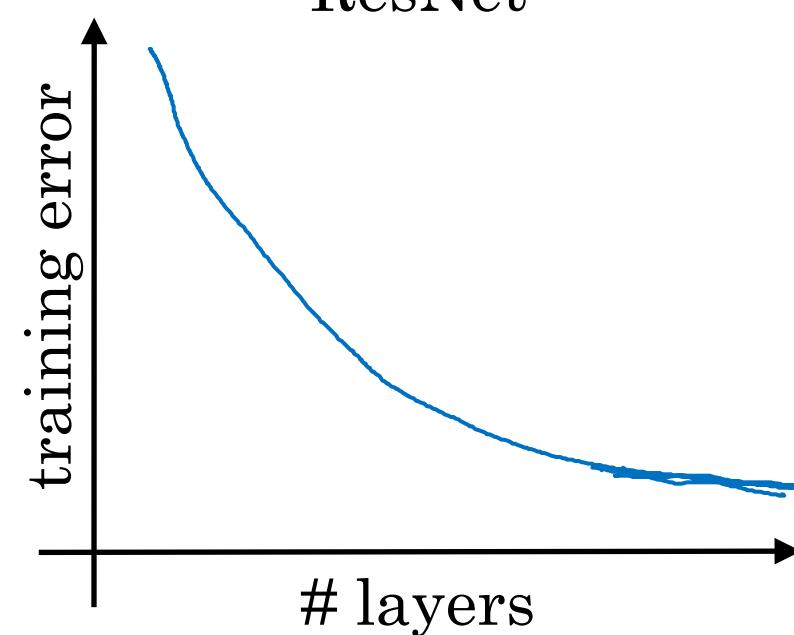
# Residual Network



Plain



ResNet





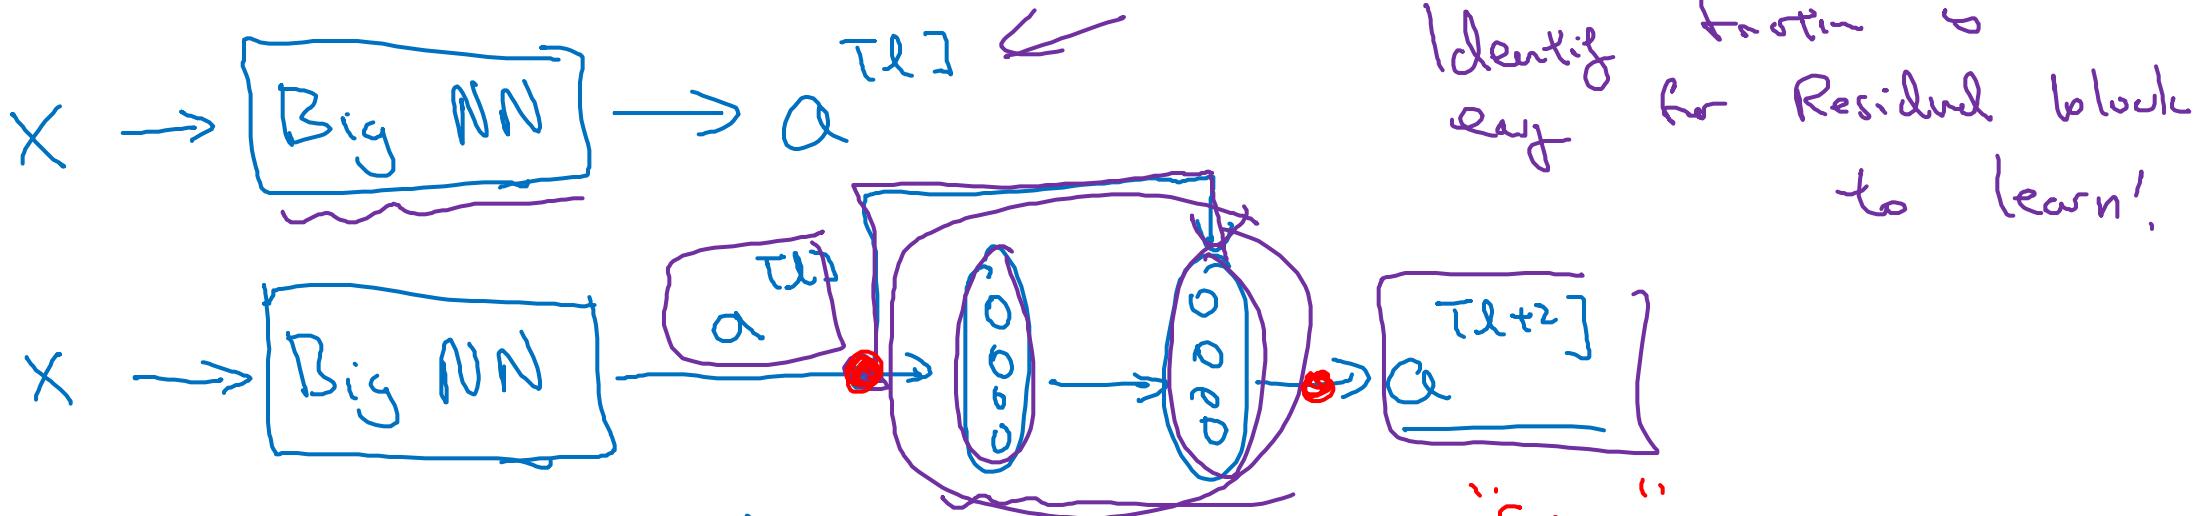
deeplearning.ai

# Case Studies

---

## Why ResNets work

# Why do residual networks work?

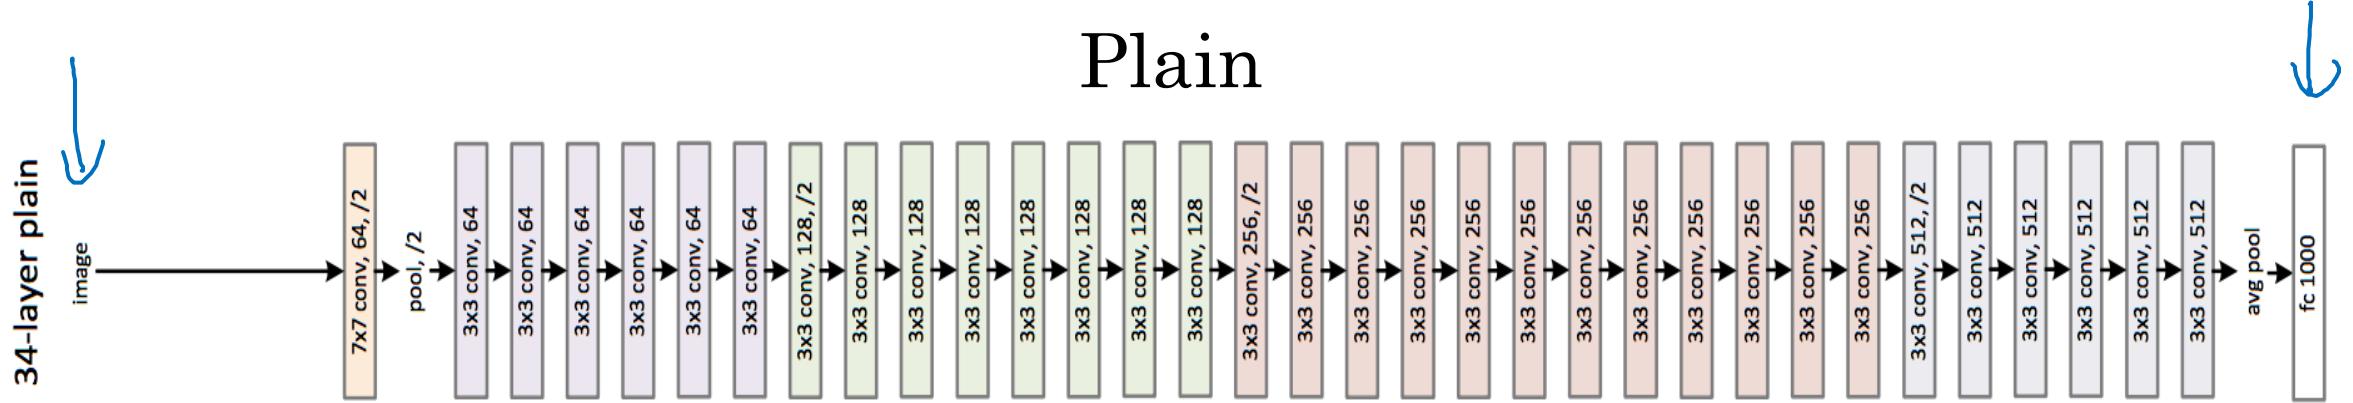


ReLU.

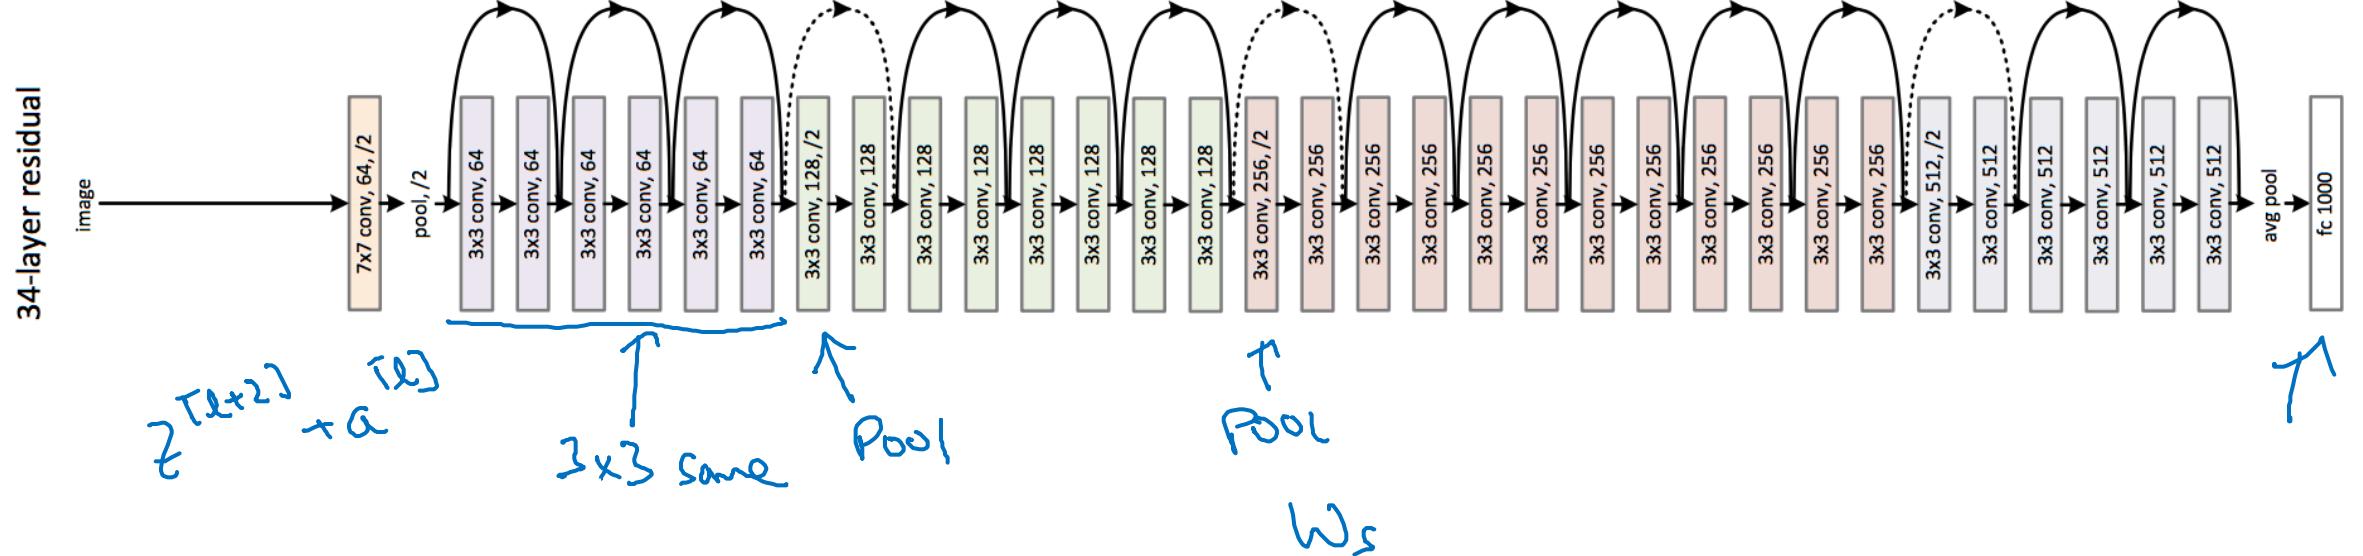
$$\begin{aligned}
 a^{[l+2]} &= g(z^{[l+2]} + a^{[l]}) \\
 &= g(w^{[l+2]} a^{[l+1]} + b^{[l+2]}) + w_s \underbrace{a^{[l]}}_{128} \\
 &\quad \text{If } w^{[l+2]} = 0, b^{[l+2]} = 0
 \end{aligned}$$

$256 \times 128$   
 $R$   
 $w_s$   
 $a^{[l]}$   
 $128$

# ResNet



Plain





deeplearning.ai

## Case Studies

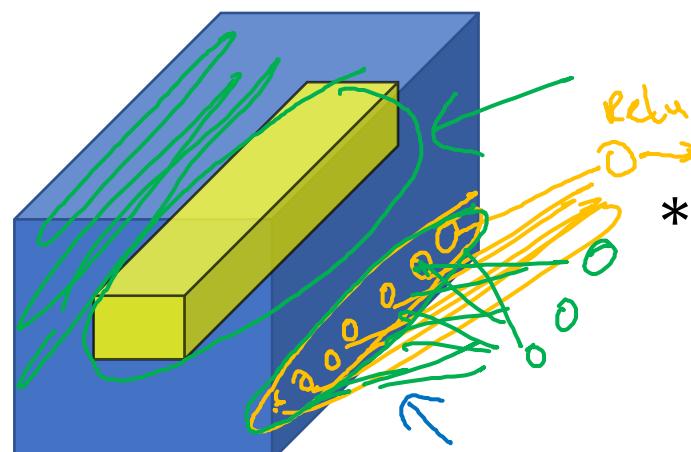
---

# Network in Network and $1 \times 1$ convolutions

# Why does a $1 \times 1$ convolution do?

1	2	3	6	5	8
3	5	5	1	3	4
2	1	3	4	9	3
4	7	8	5	7	9
1	5	3	7	4	8
5	4	9	8	3	5

$6 \times 6$



$6 \times 6 \times 32$

[Lin et al., 2013. Network in network]

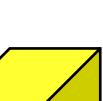
\*

2



=

32



# filters.  
 $n_c^{[l+1]}$

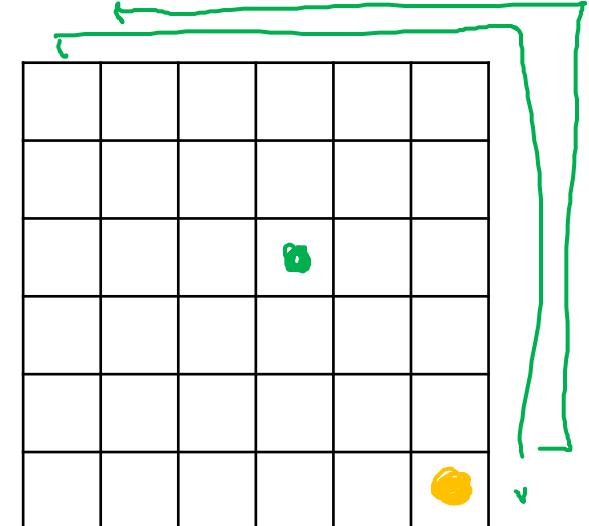
=

ReLU

Network  $\rightarrow$  Network

$1 \times 1 \times 32$

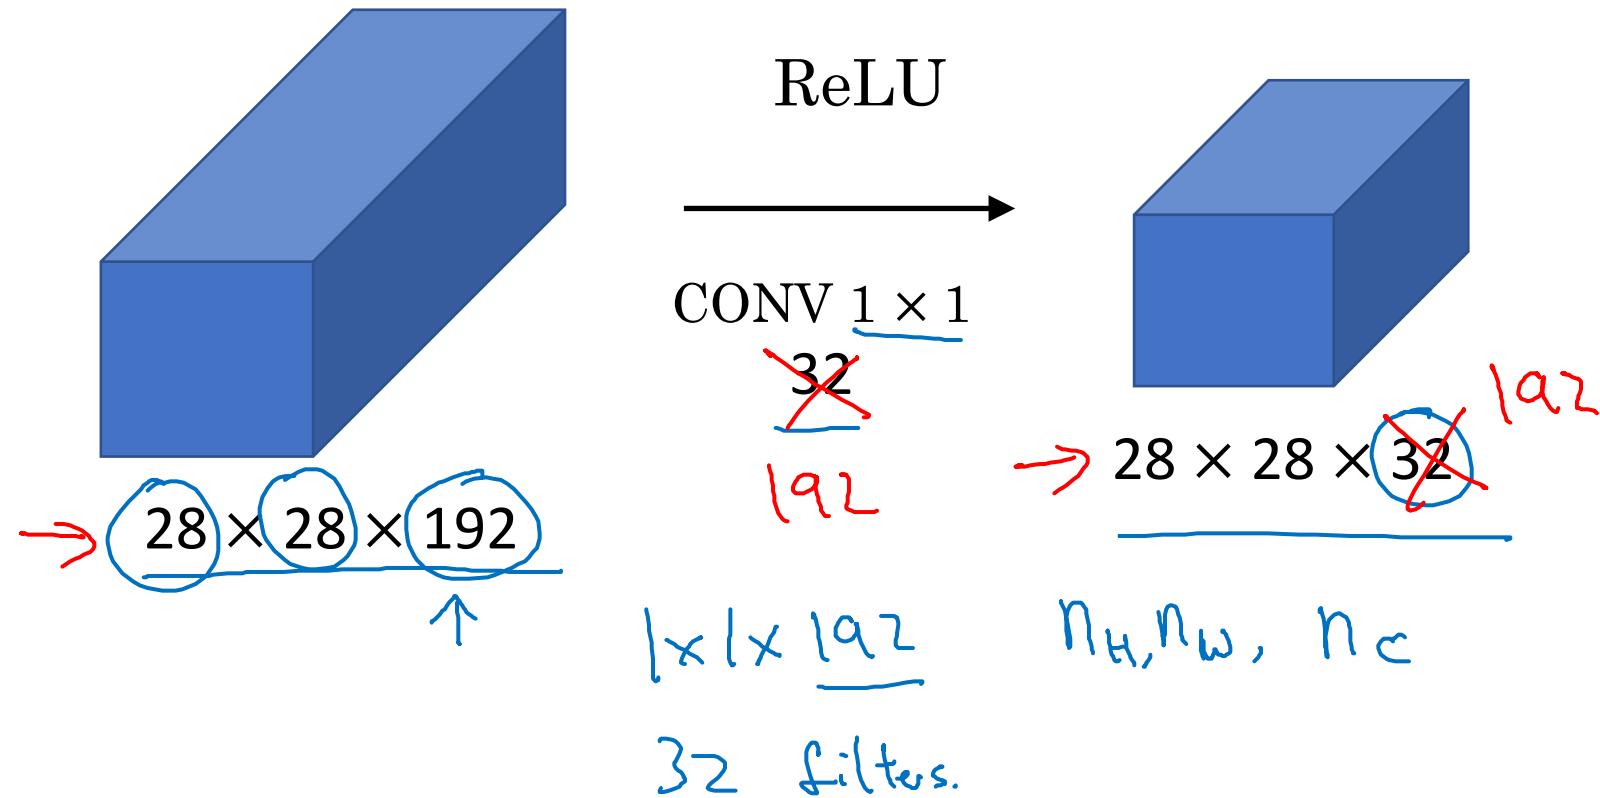
2	4	6	...



$6 \times 6 \times \# \text{ filters}$

Andrew Ng

# Using $1 \times 1$ convolutions





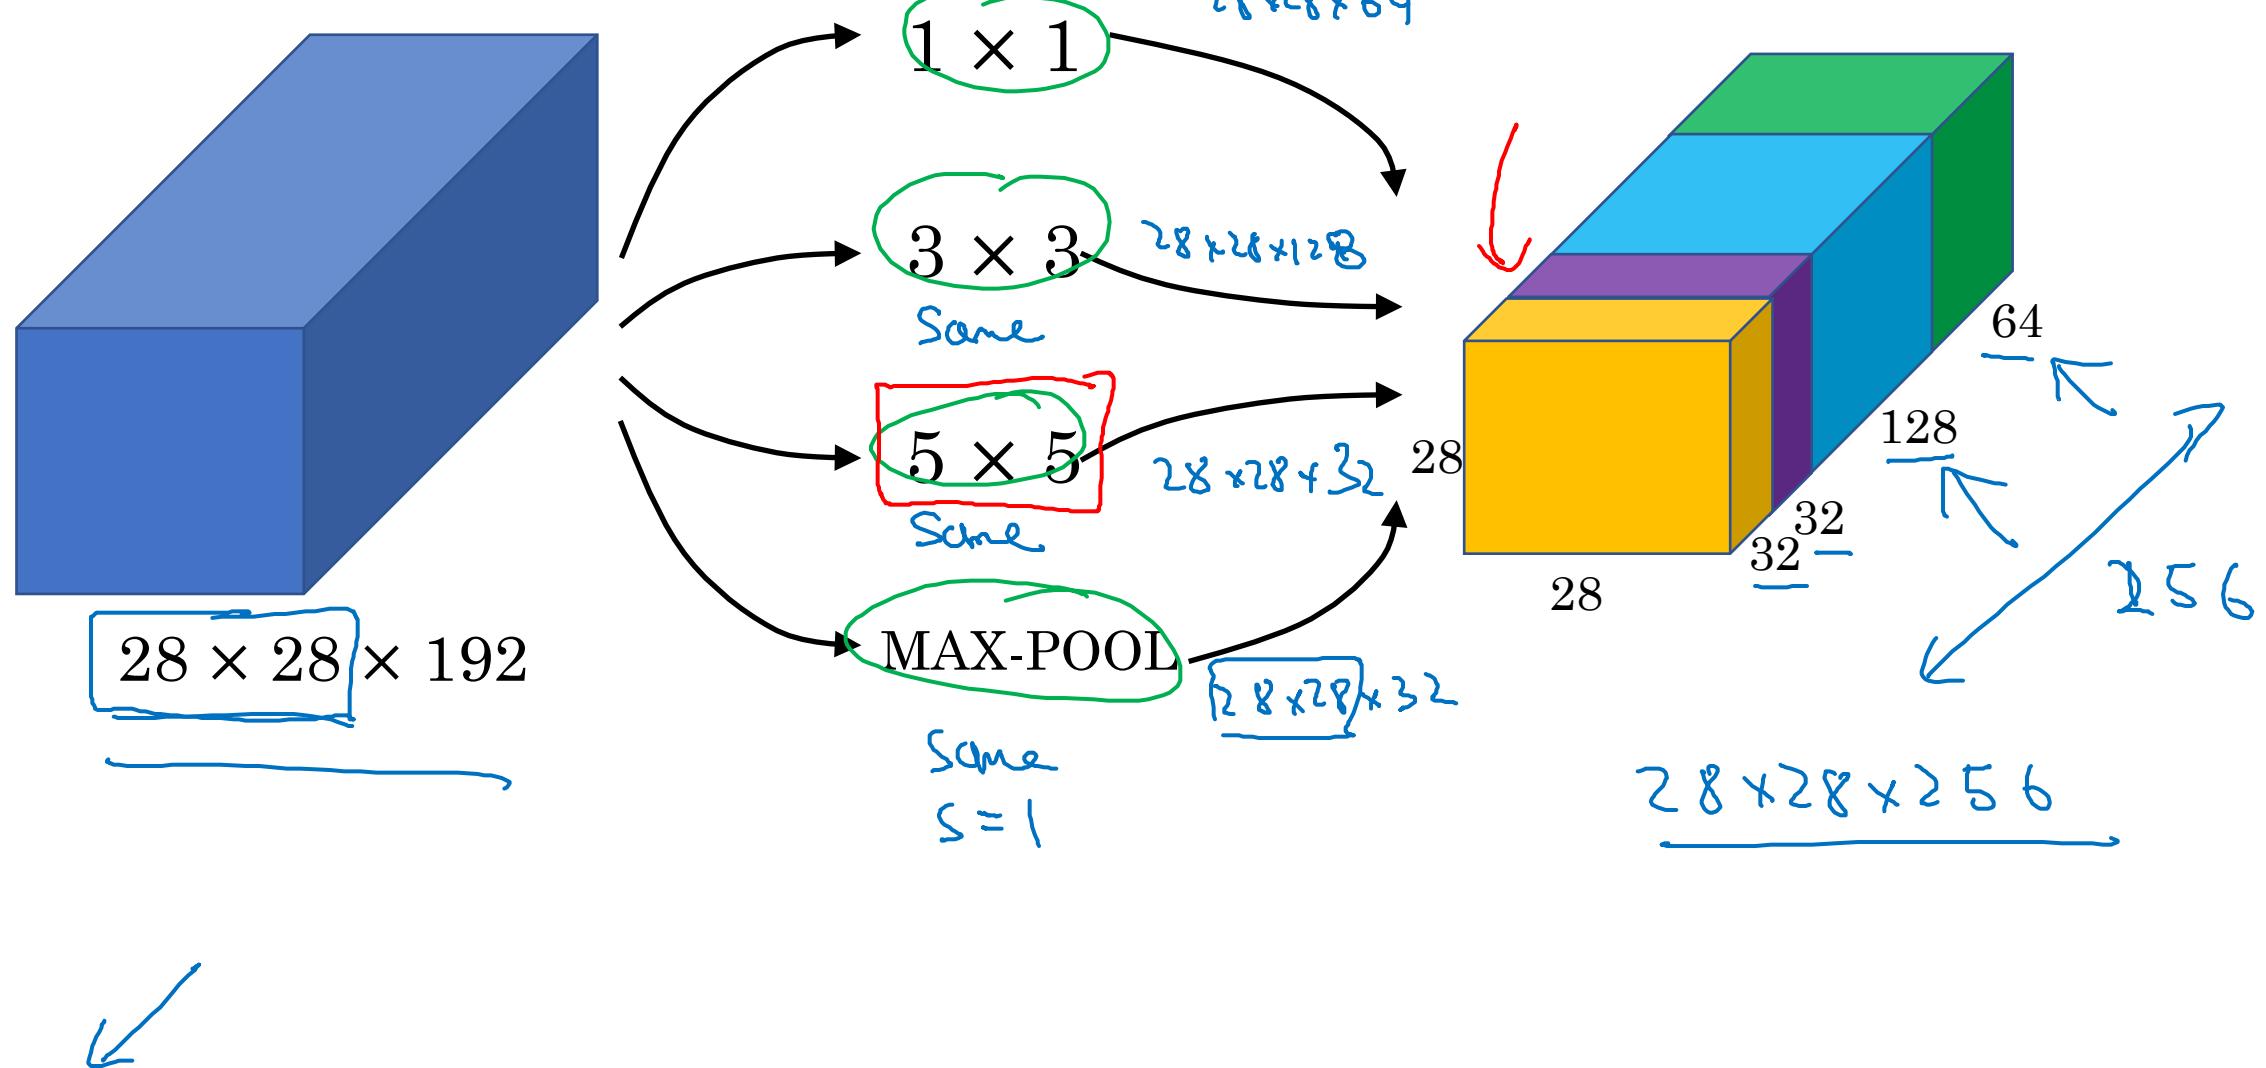
deeplearning.ai

## Case Studies

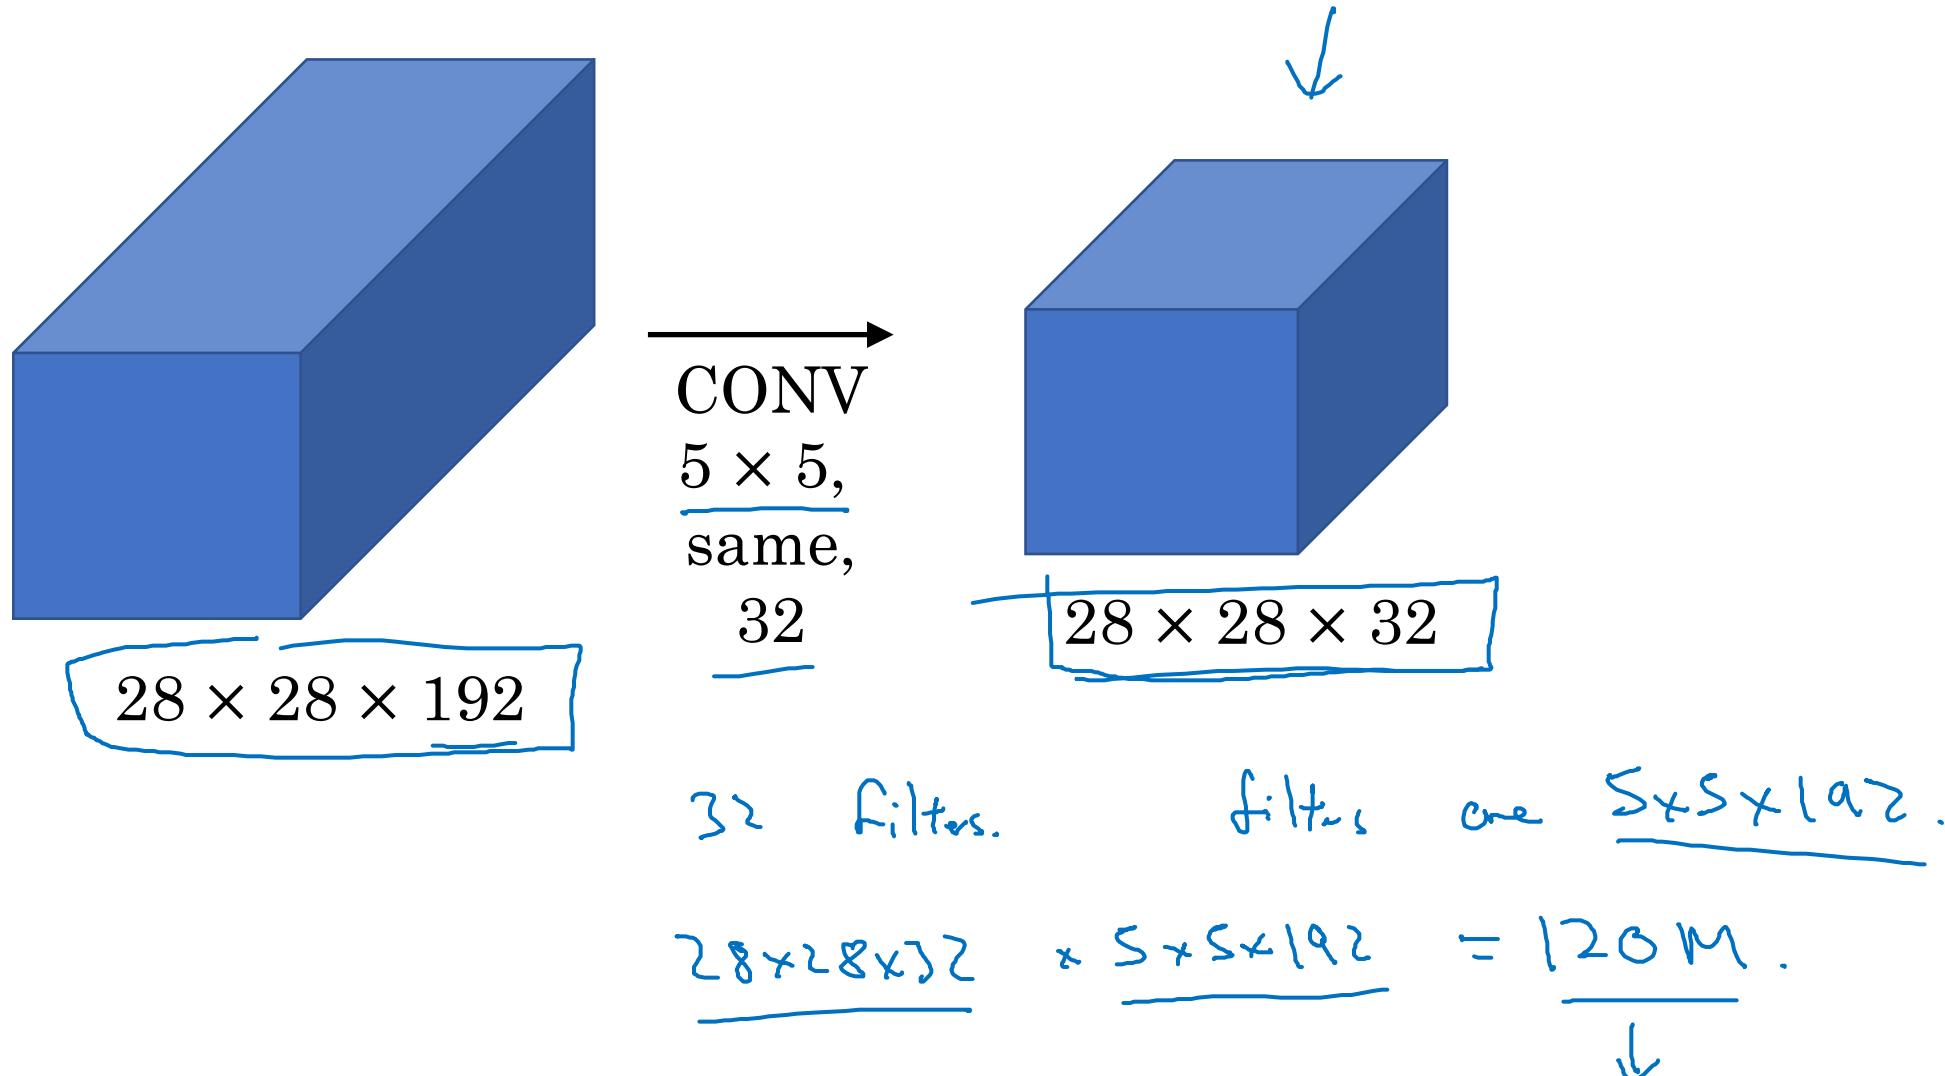
---

Inception network  
motivation

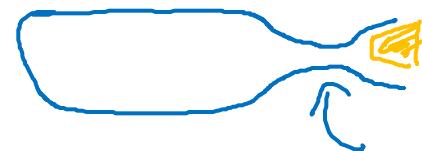
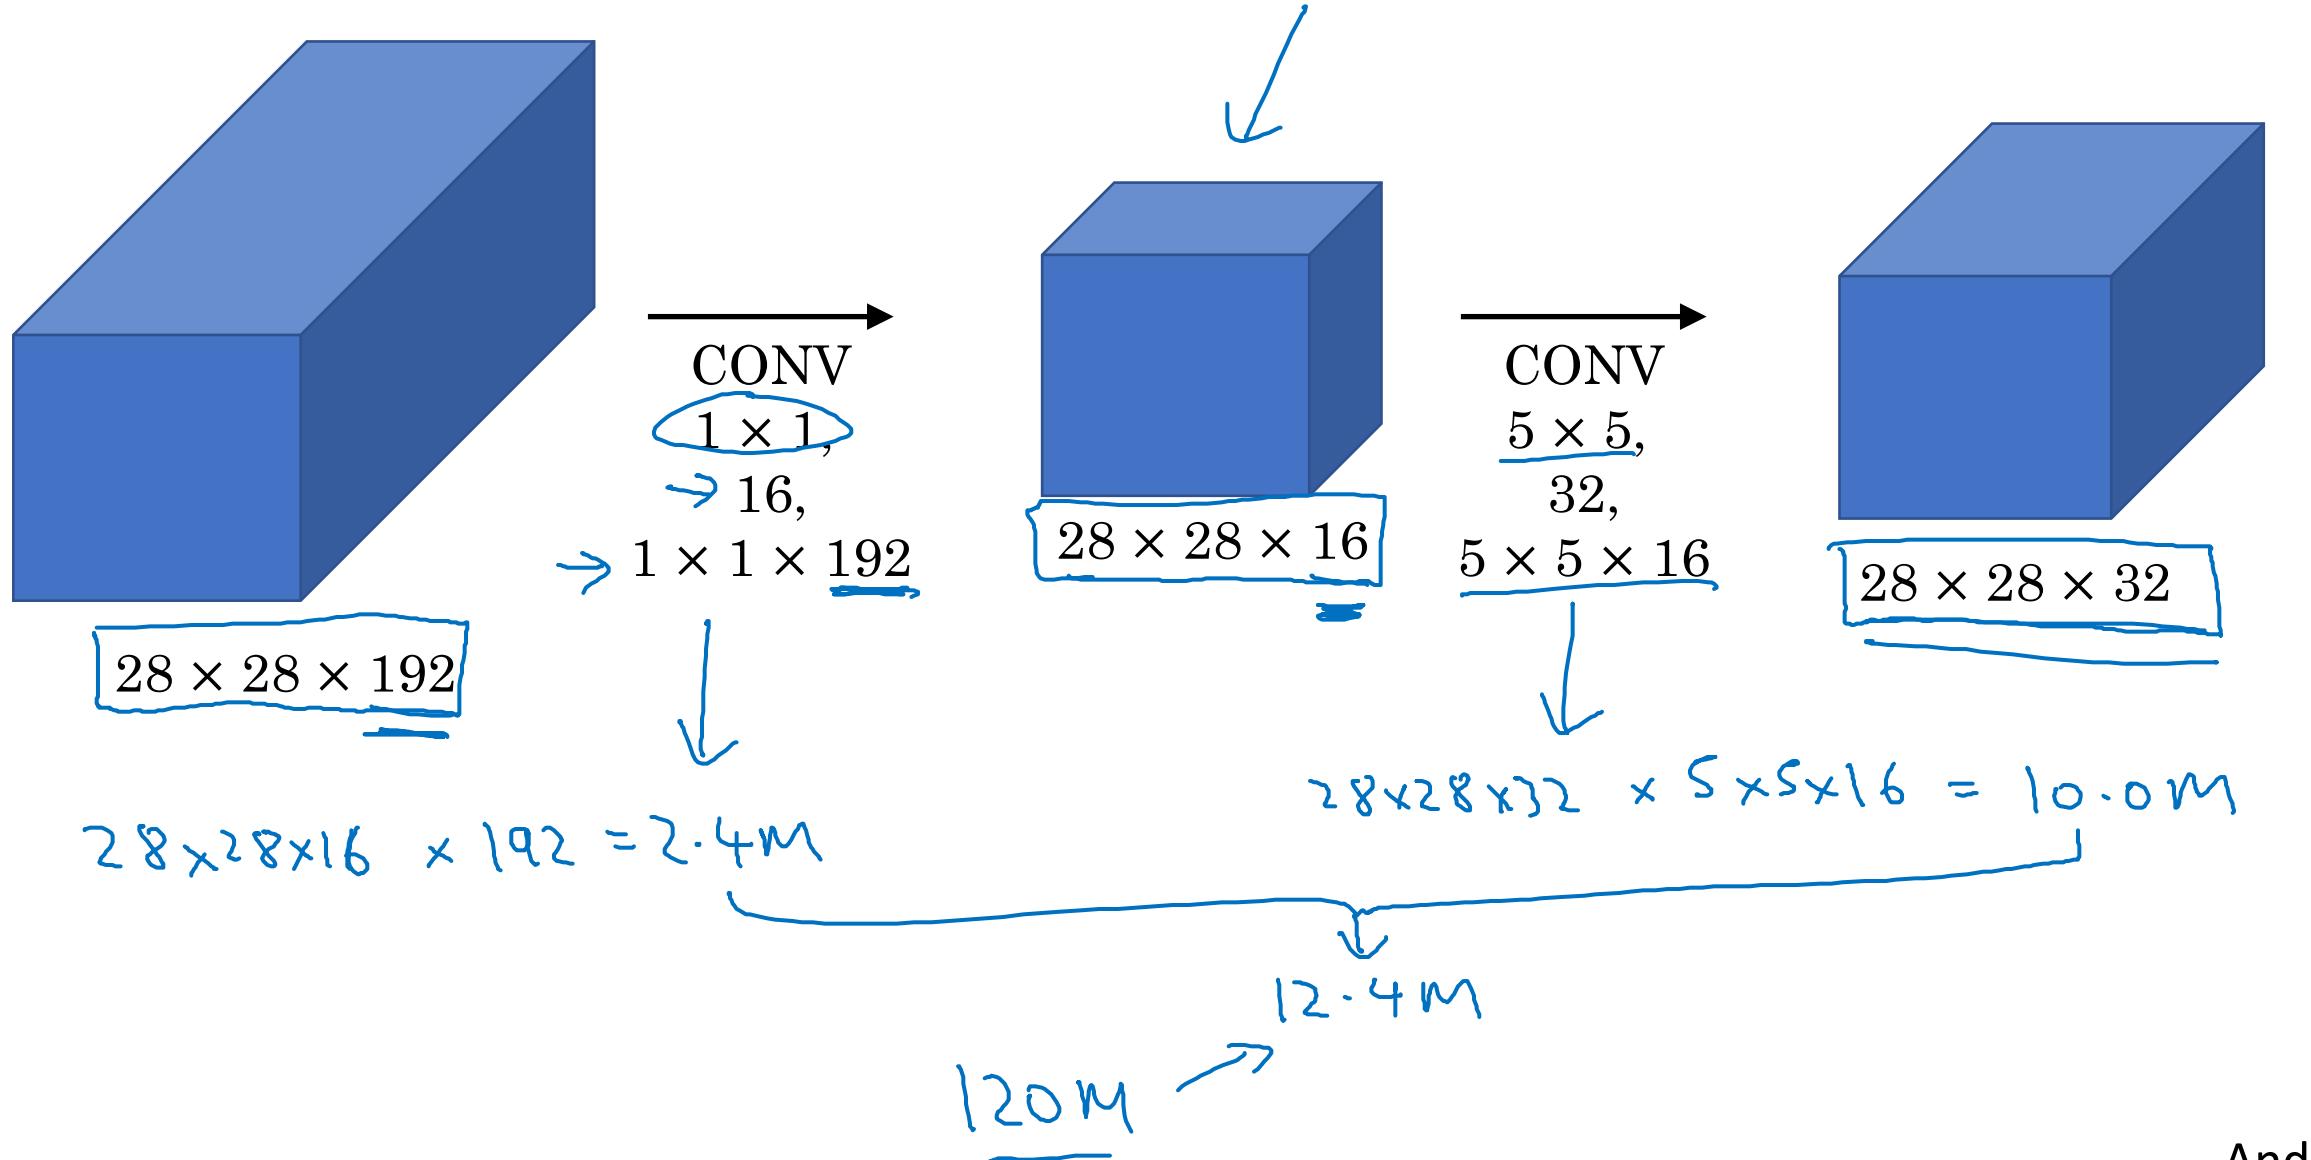
# Motivation for inception network



# The problem of computational cost



# Using $1 \times 1$ convolution





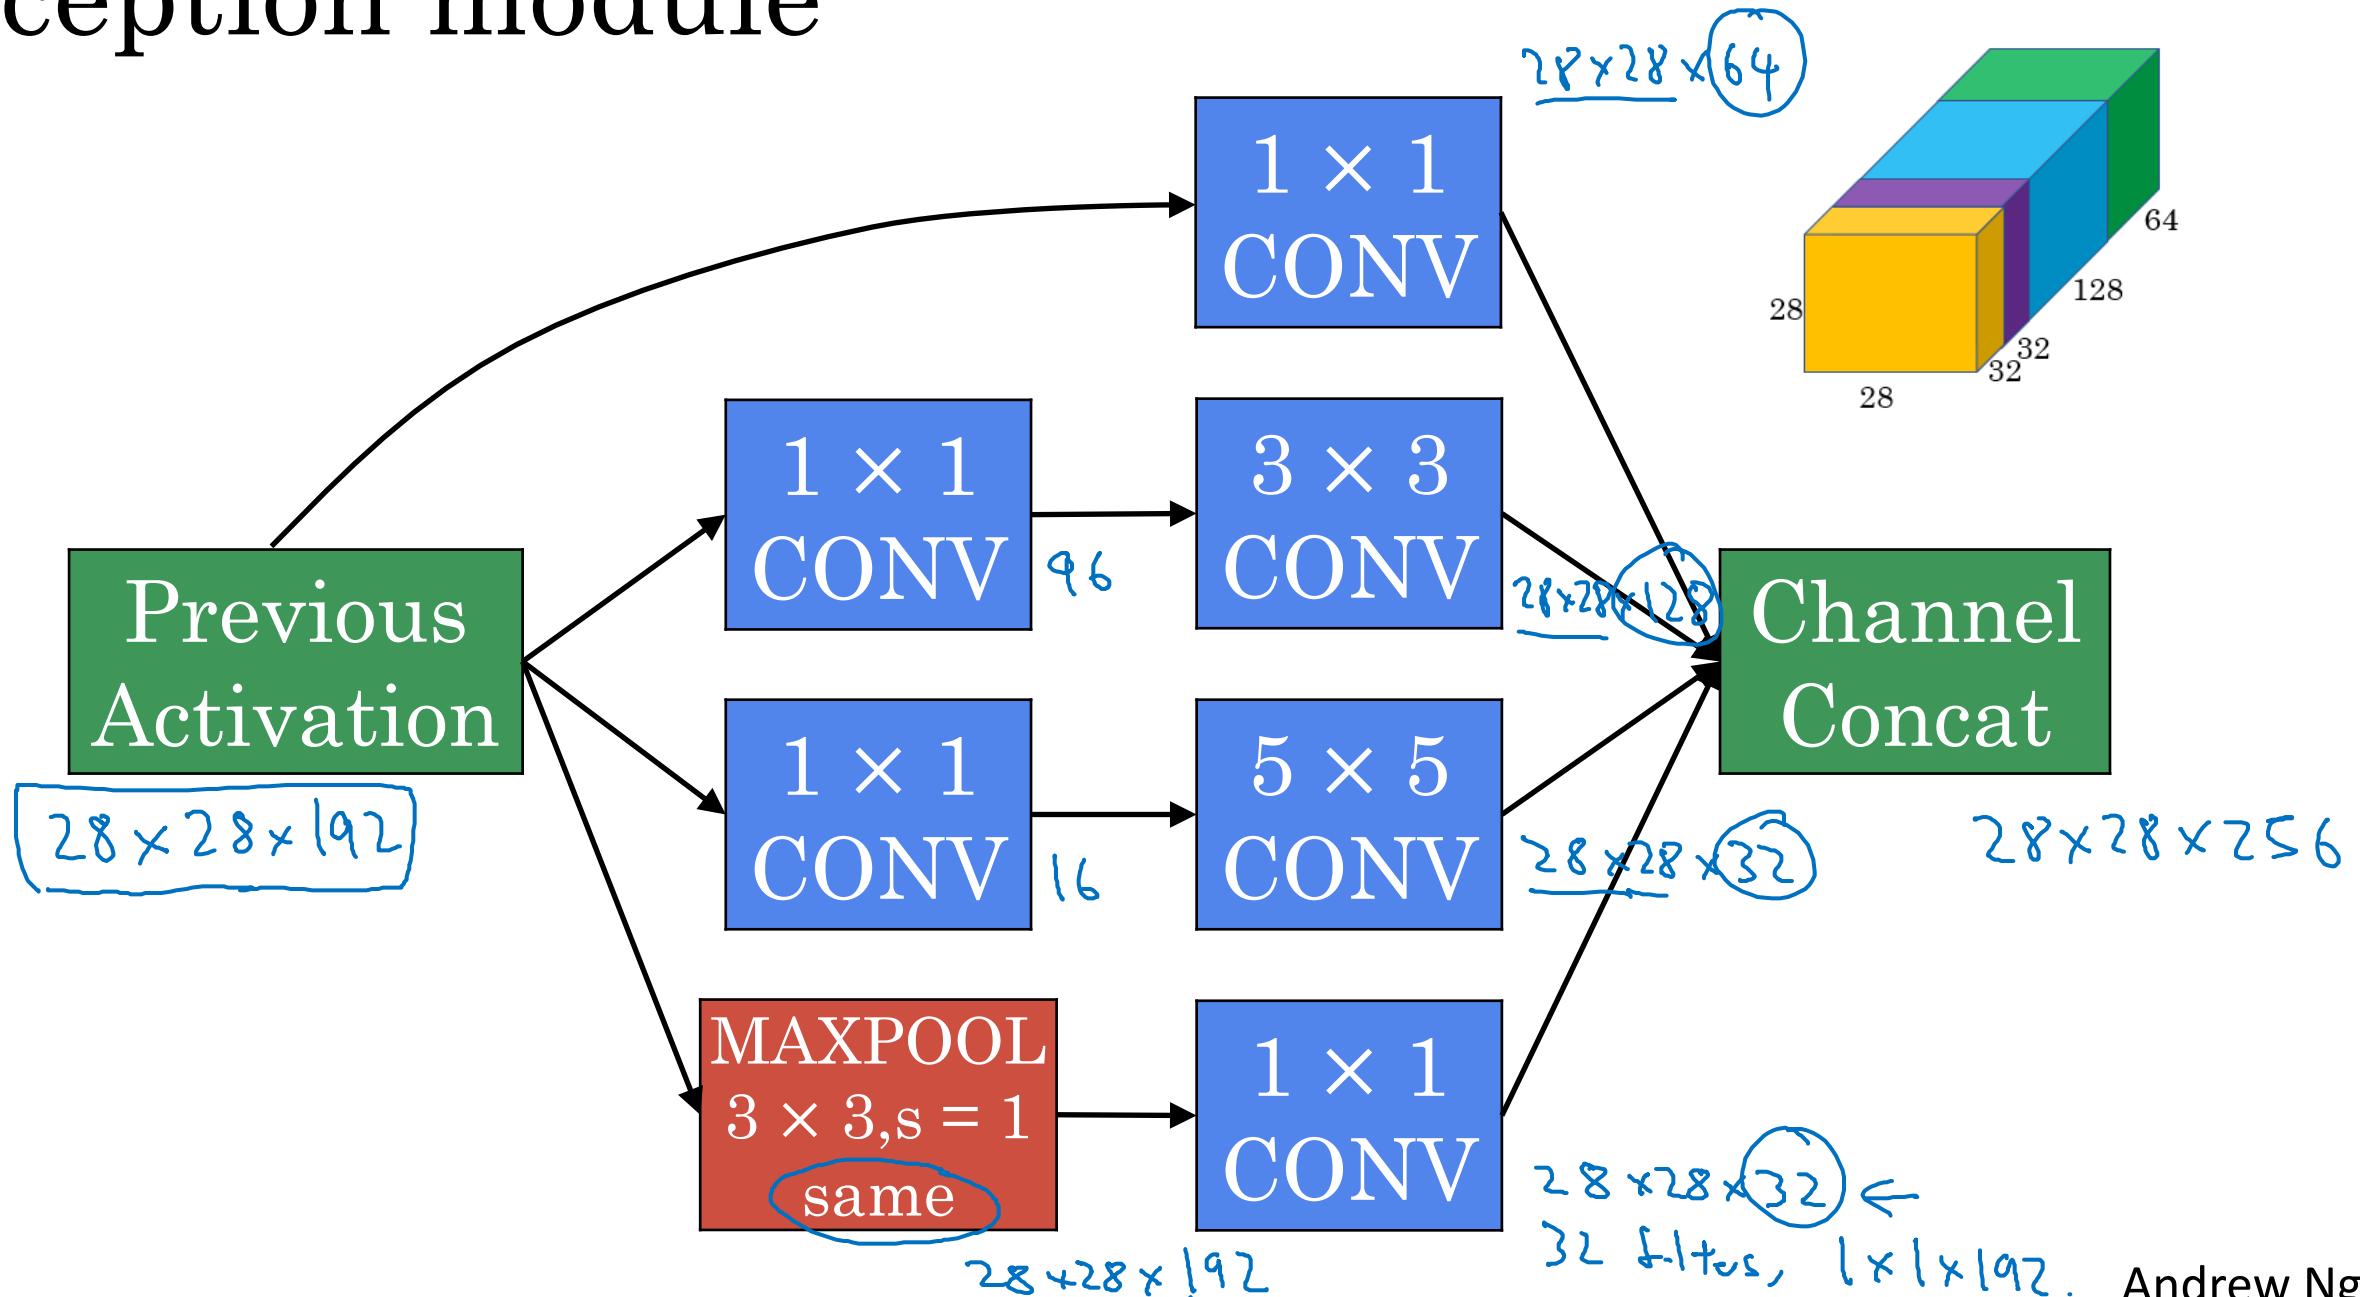
deeplearning.ai

## Case Studies

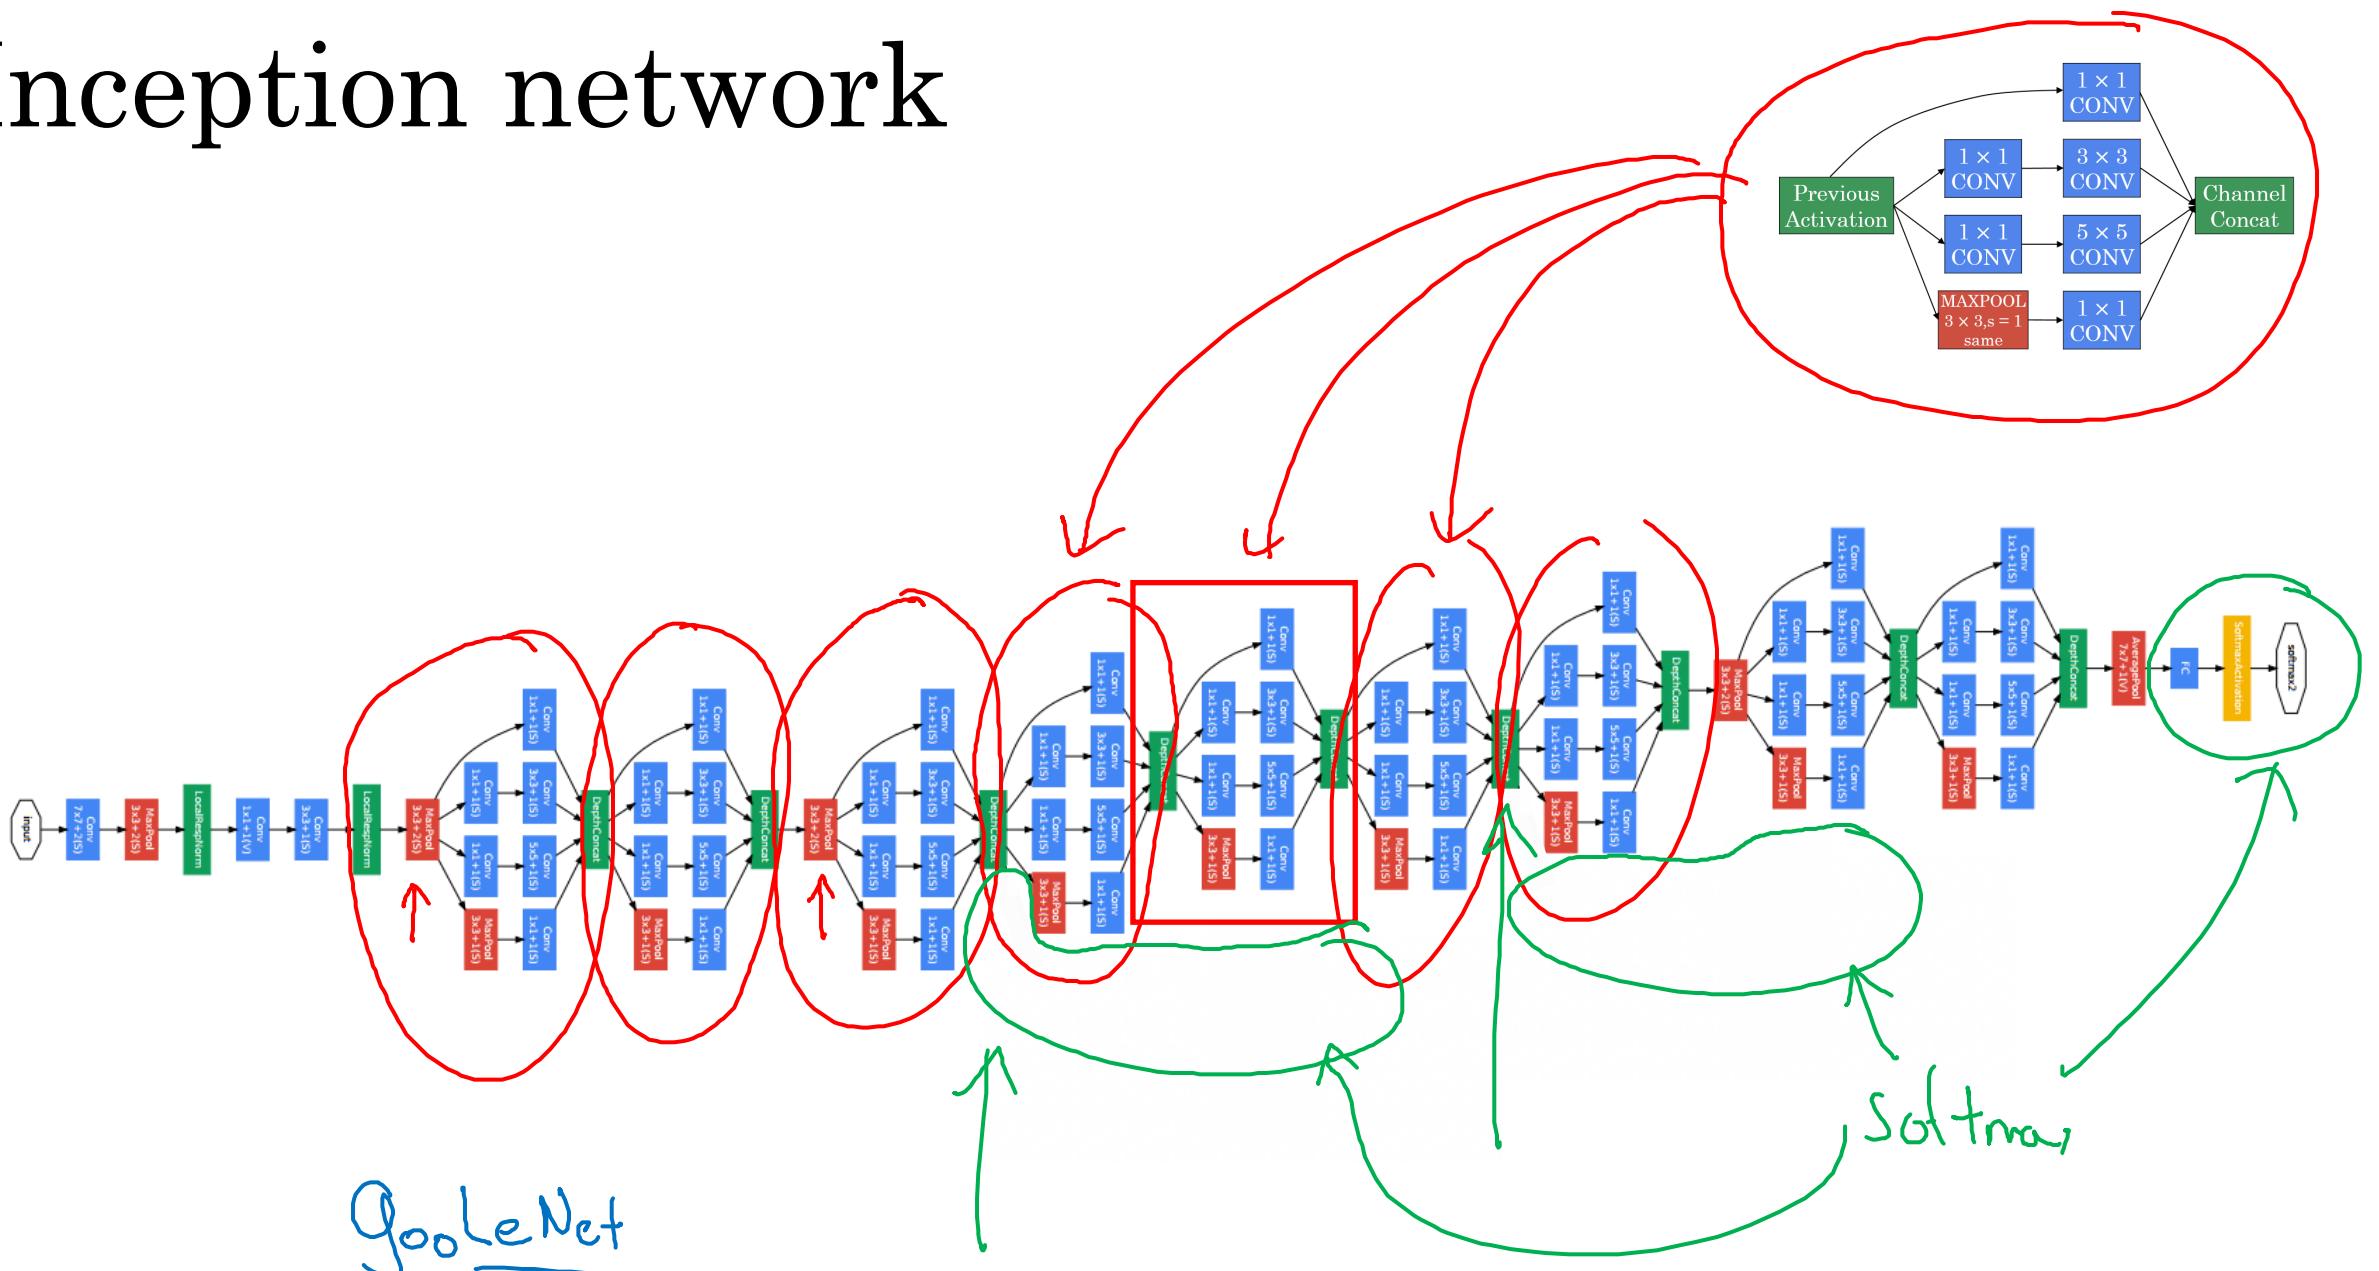
---

### Inception network

# Inception module



# Inception network







deeplearning.ai

# Convolutional Neural Networks

---

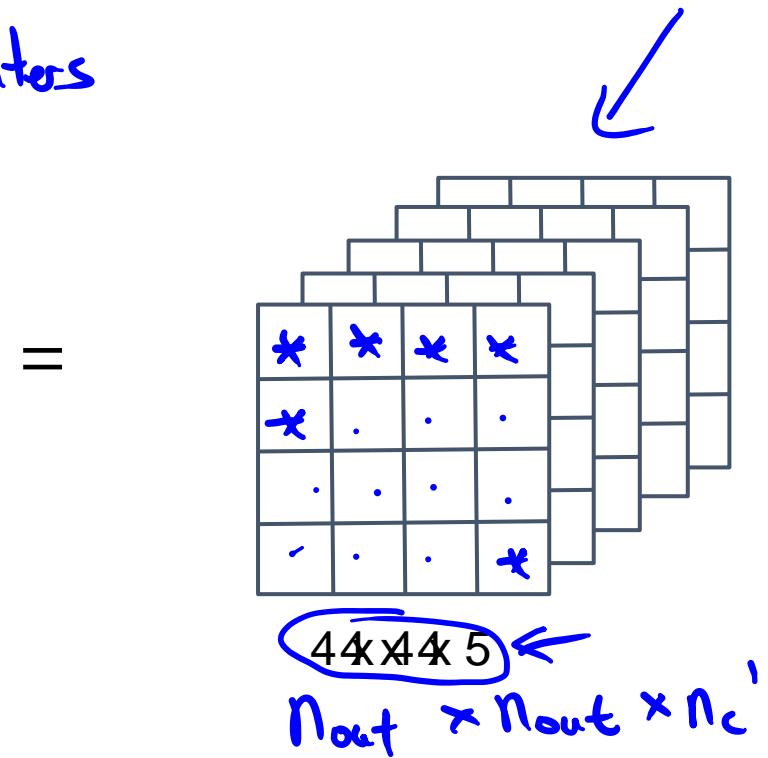
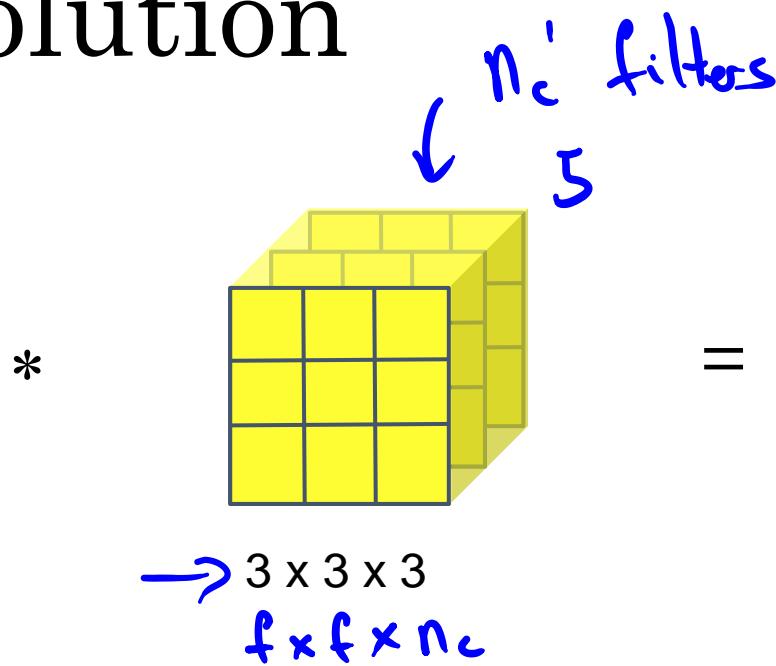
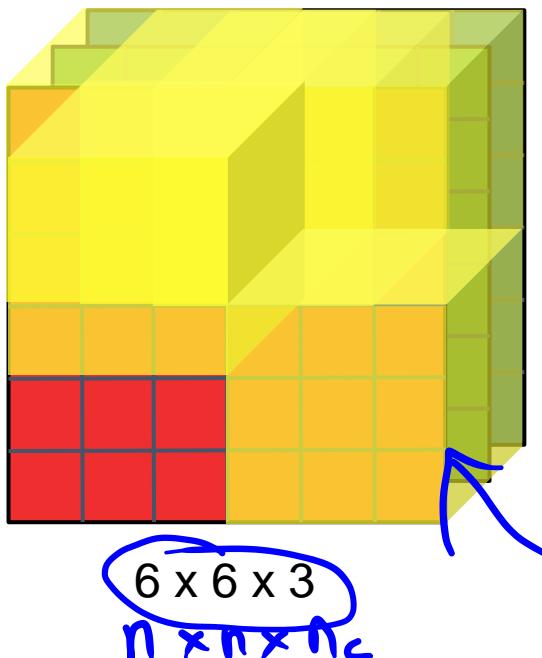
## MobileNet

# Motivation for MobileNets

- Low computational cost at deployment
- Useful for mobile and embedded vision applications
- Key idea: Normal vs. depthwise-separable convolutions



# Normal Convolution



$$\text{Computational cost} = \frac{\# \text{filter params}}{3 \times 3 \times 3} \times \frac{\# \text{filter positions}}{4 \times 4}$$

$\rightarrow \underline{2160} =$

↑                      ↑

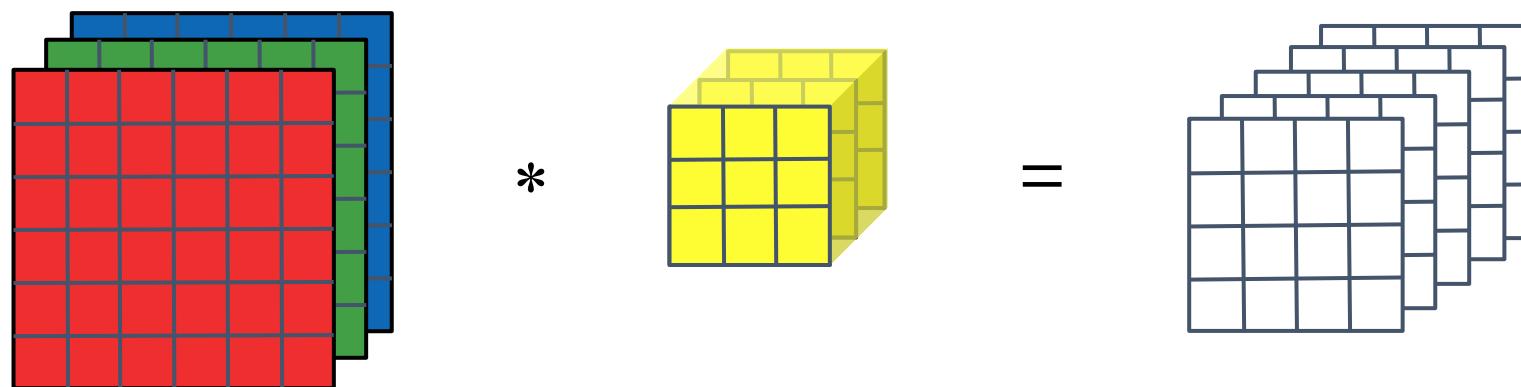
$$x \quad \# \text{of filters}$$

$\times \quad 5$

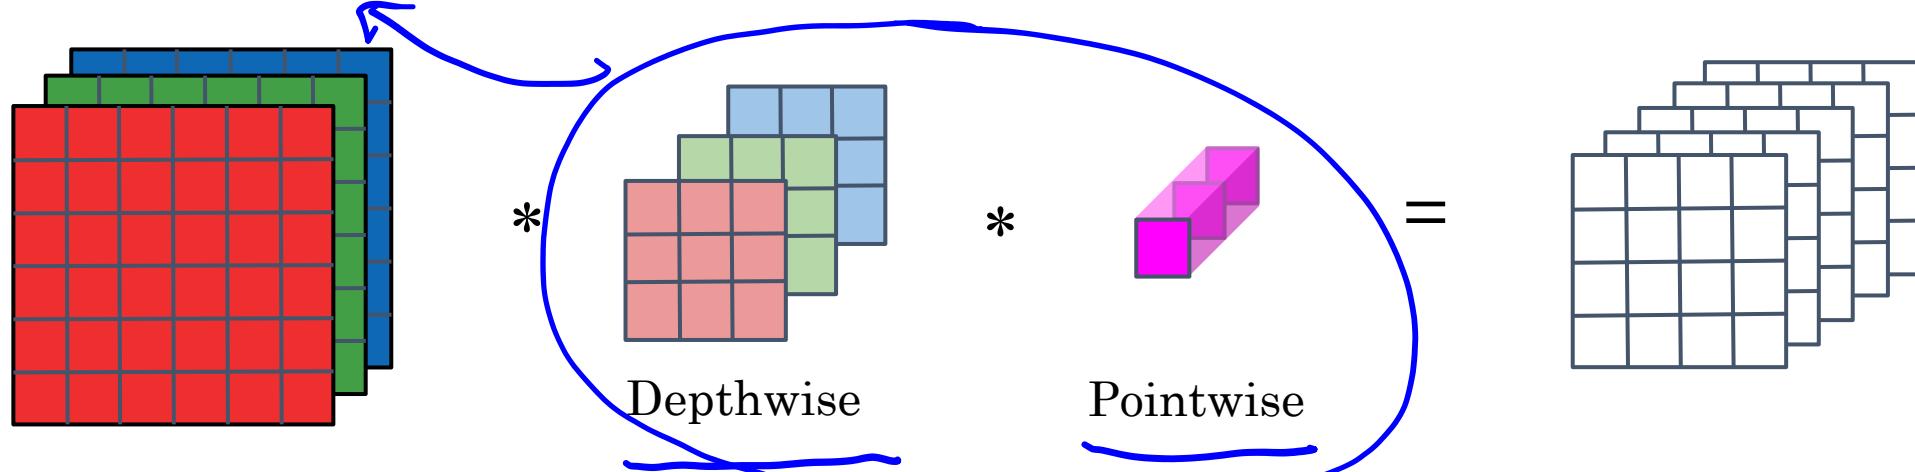
↑

# Depthwise Separable Convolution

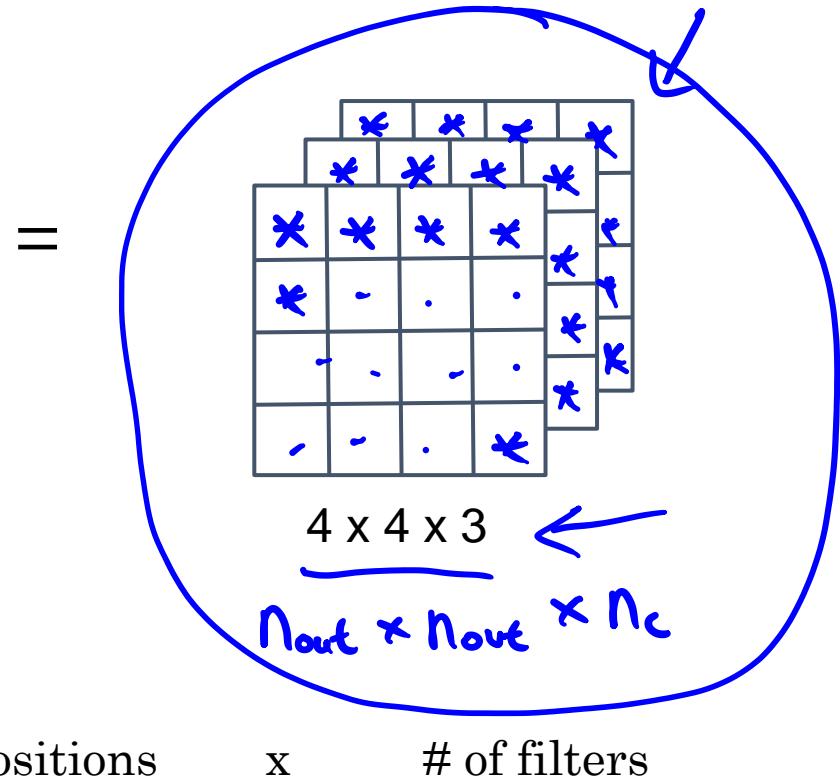
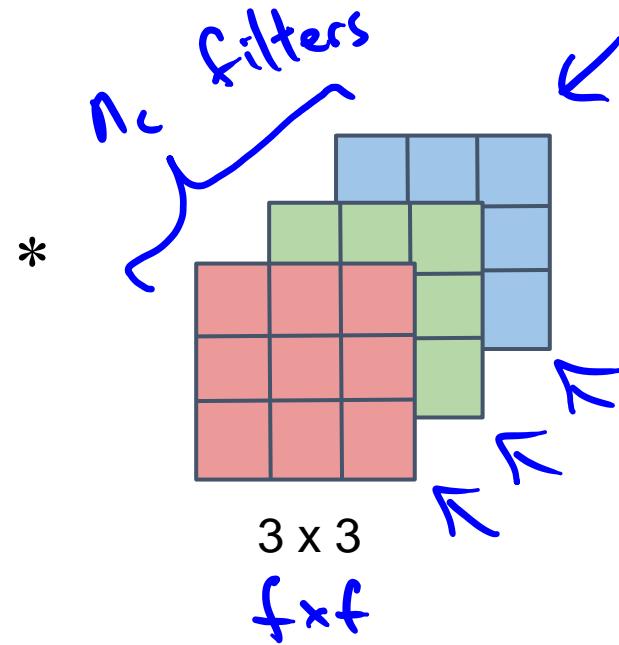
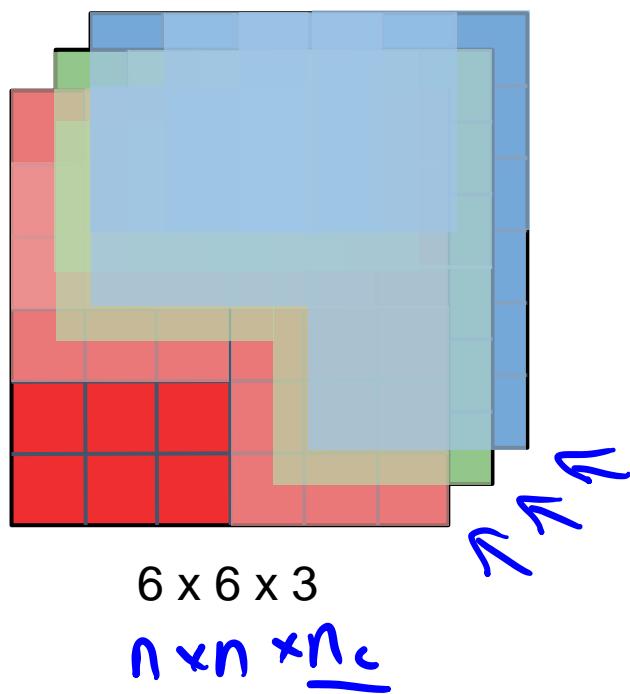
Normal Convolution



Depthwise Separable Convolution



# Depthwise Convolution

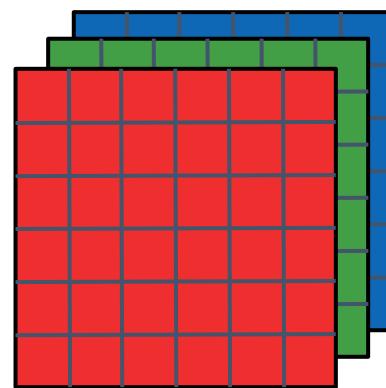


Computational cost = #filter params  $\times$  # filter positions  $\times$  # of filters

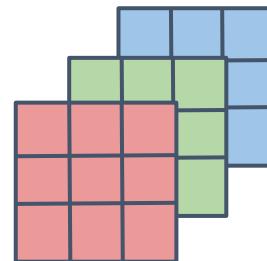
$$432 = \underbrace{3 \times 3}_{\# \text{filter params}} \times \underbrace{4 \times 4}_{\# \text{filter positions}} \times \underbrace{3}_{\# \text{of filters}}$$

# Depthwise Separable Convolution

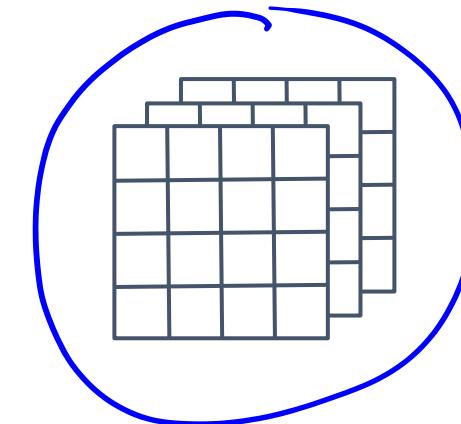
Depthwise Convolution



\*

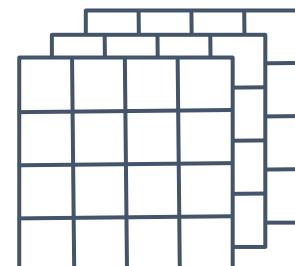


=



432

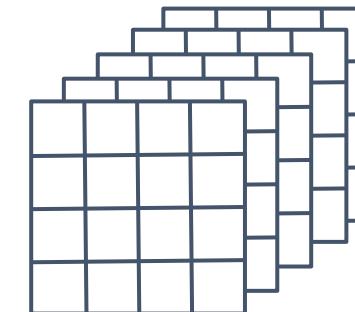
Pointwise Convolution



\*

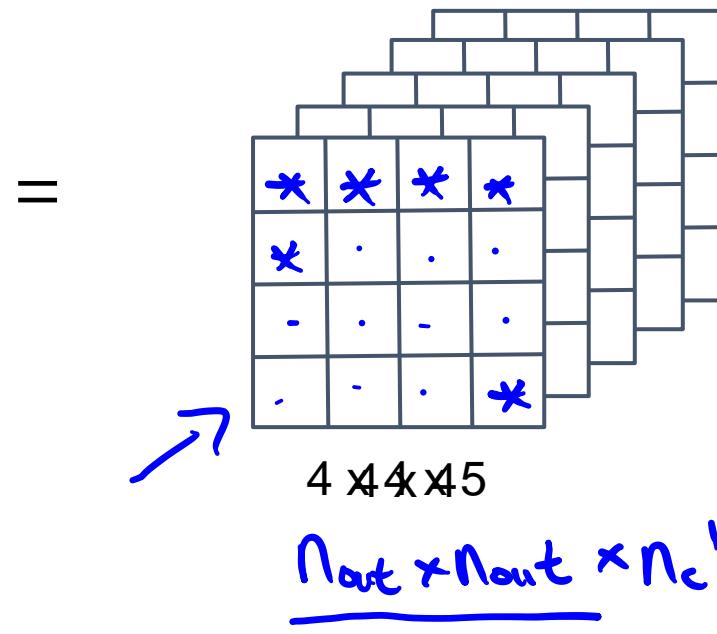
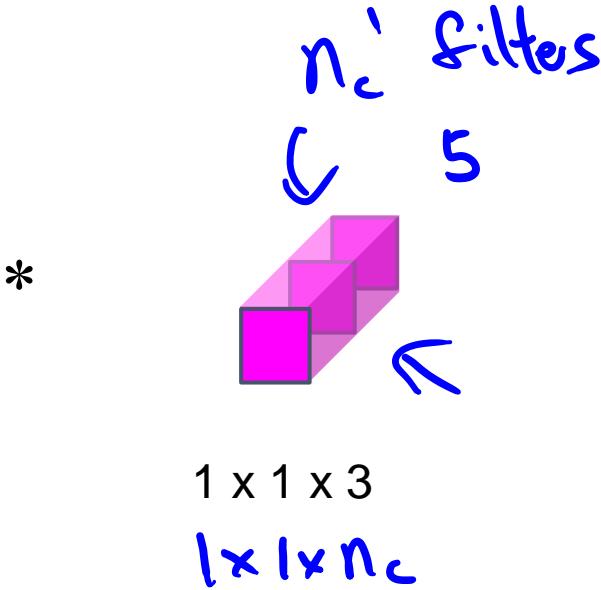
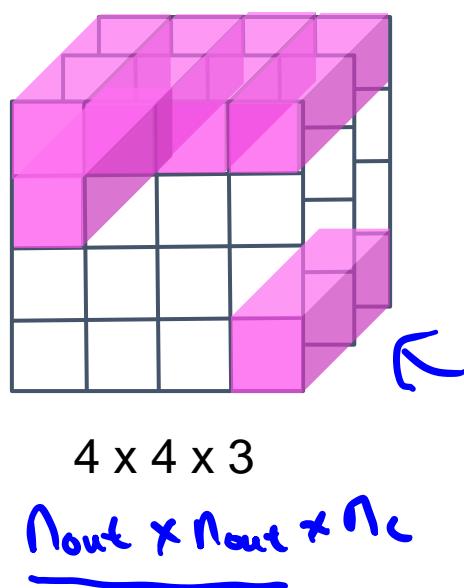


=



4x4x5

# Pointwise Convolution

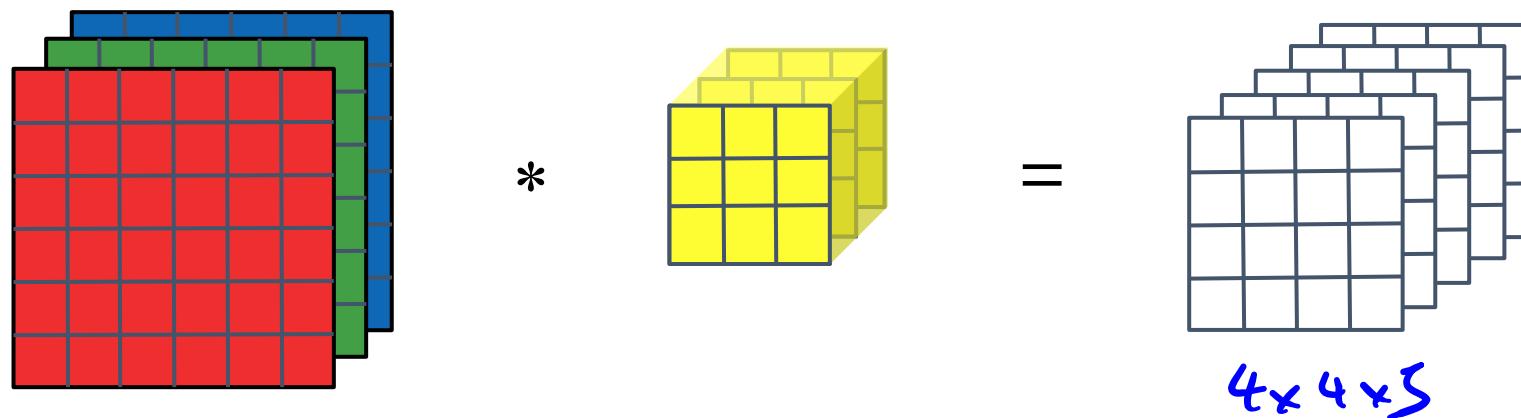


Computational cost = #filter params  $\times$  # filter positions  $\times$  # of filters

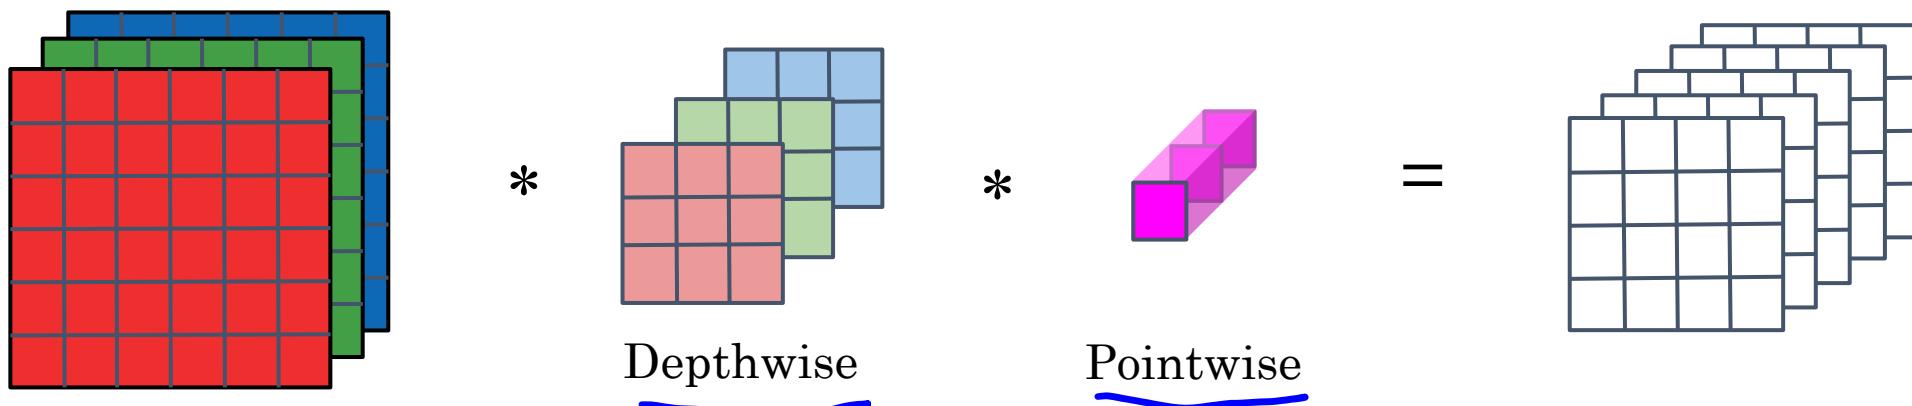
$$240 = 1 \times 1 \times 3 \times 4 \times 4 \times 5$$

# Depthwise Separable Convolution

Normal Convolution



Depthwise Separable Convolution



# Cost Summary

Cost of normal convolution

2160

Cost of depthwise separable convolution

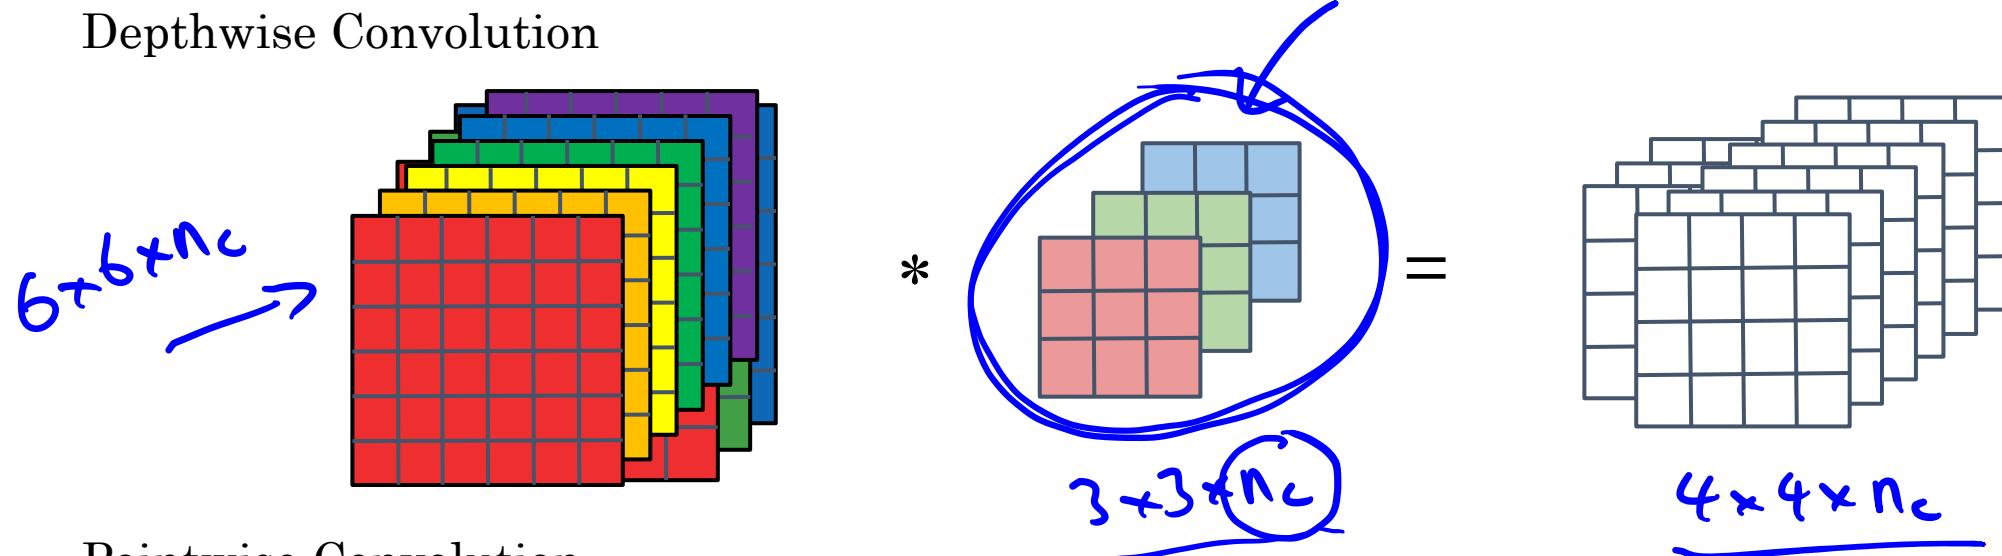
$$\begin{array}{l} \text{depthwise} + \text{pointwise} \\ 432 + 240 = 672 \end{array}$$

$$\frac{672}{2160} = 0.31 <$$

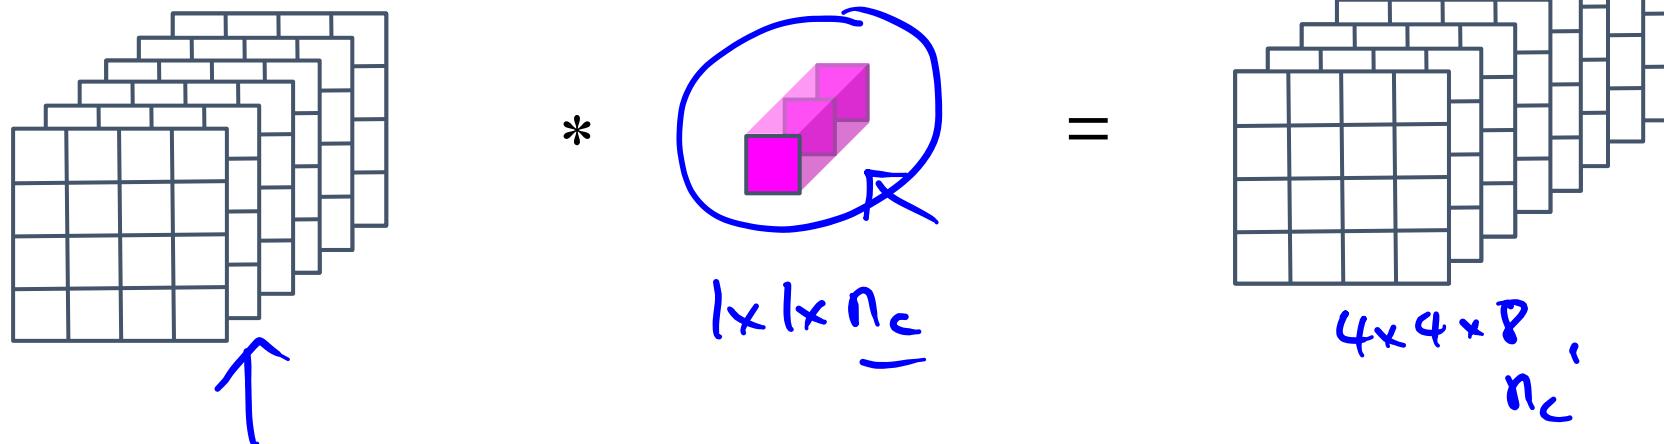
$$\begin{aligned} &= \frac{1}{n_c} + \frac{1}{f^2} \\ &\quad \frac{1}{s} + \frac{1}{q} \\ &= \frac{1}{512} + \frac{1}{3^2} \\ &\quad \nwarrow \quad \swarrow \\ &\text{n10 times cheaper} \end{aligned}$$

# Depthwise Separable Convolution

Depthwise Convolution



Pointwise Convolution





deeplearning.ai

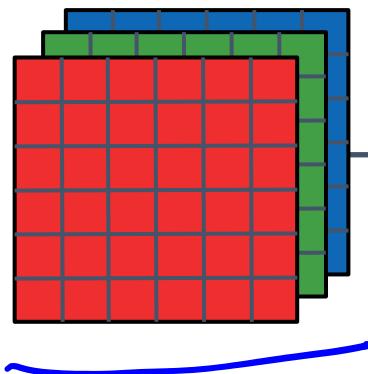
# Convolutional Neural Networks

---

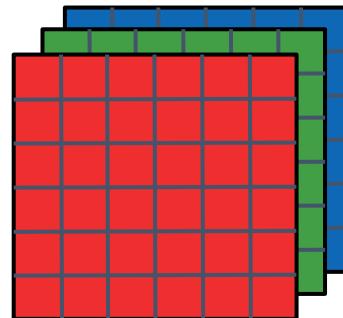
## MobileNet Architecture

# MobileNet

MobileNet v1



MobileNet v2



13 times

POOL, FC, SOFTMAX

17 times

Residual Connection

POOL, FC,  
SOFTMAX

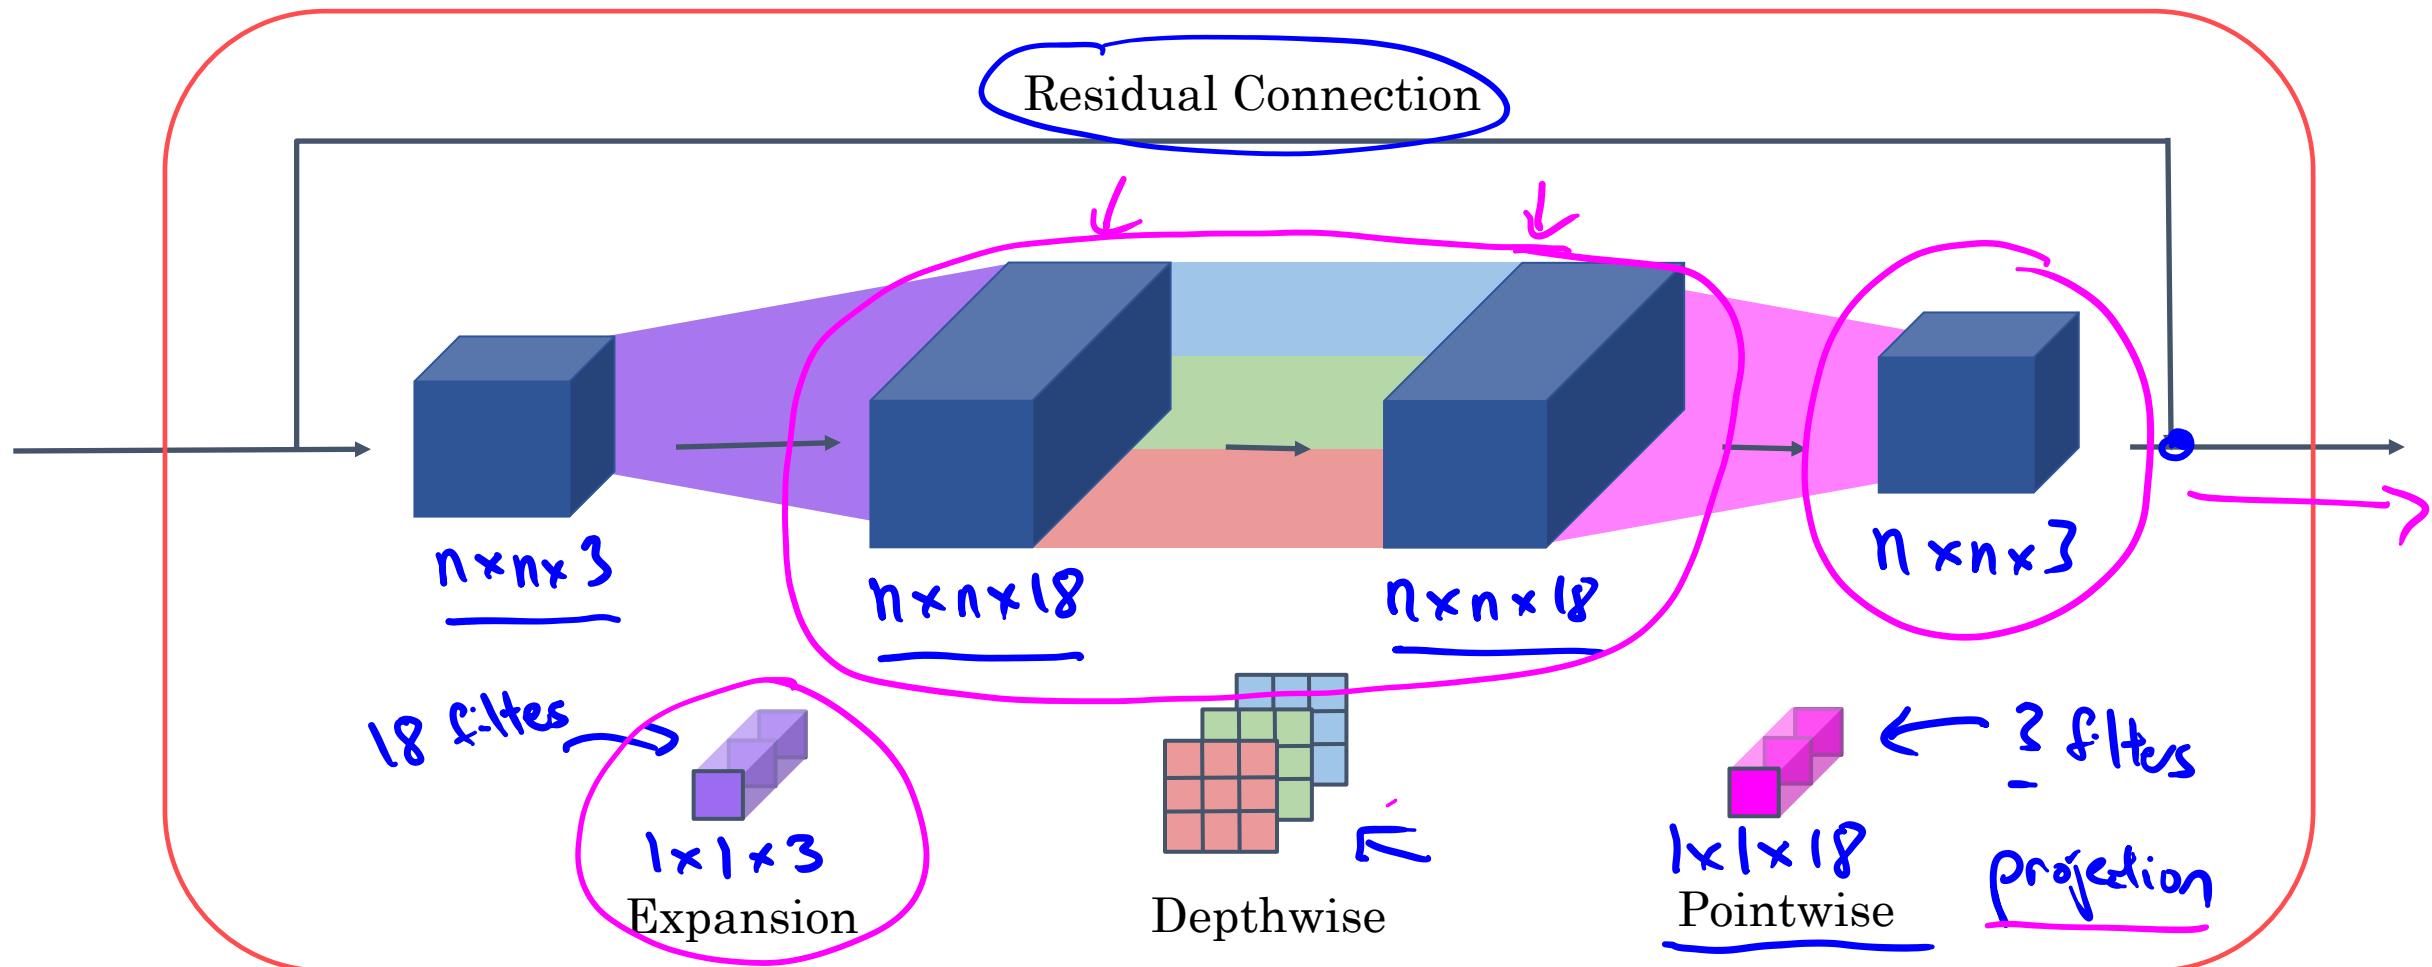
Expansion

Depthwise

Projection

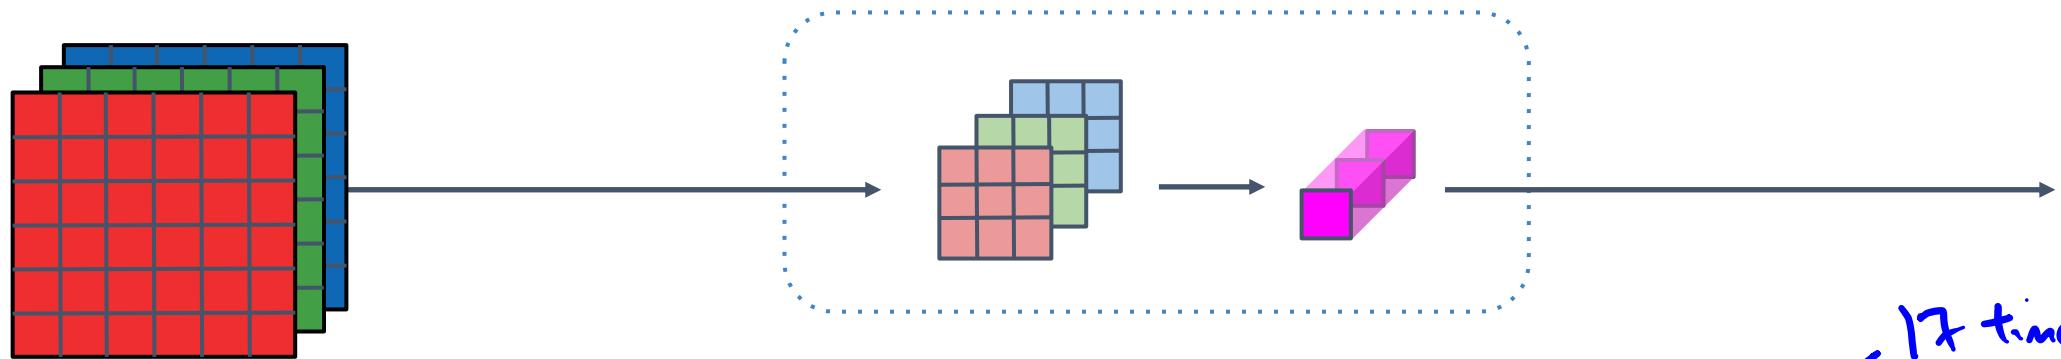
Bottleneck

# MobileNet v2 Bottleneck

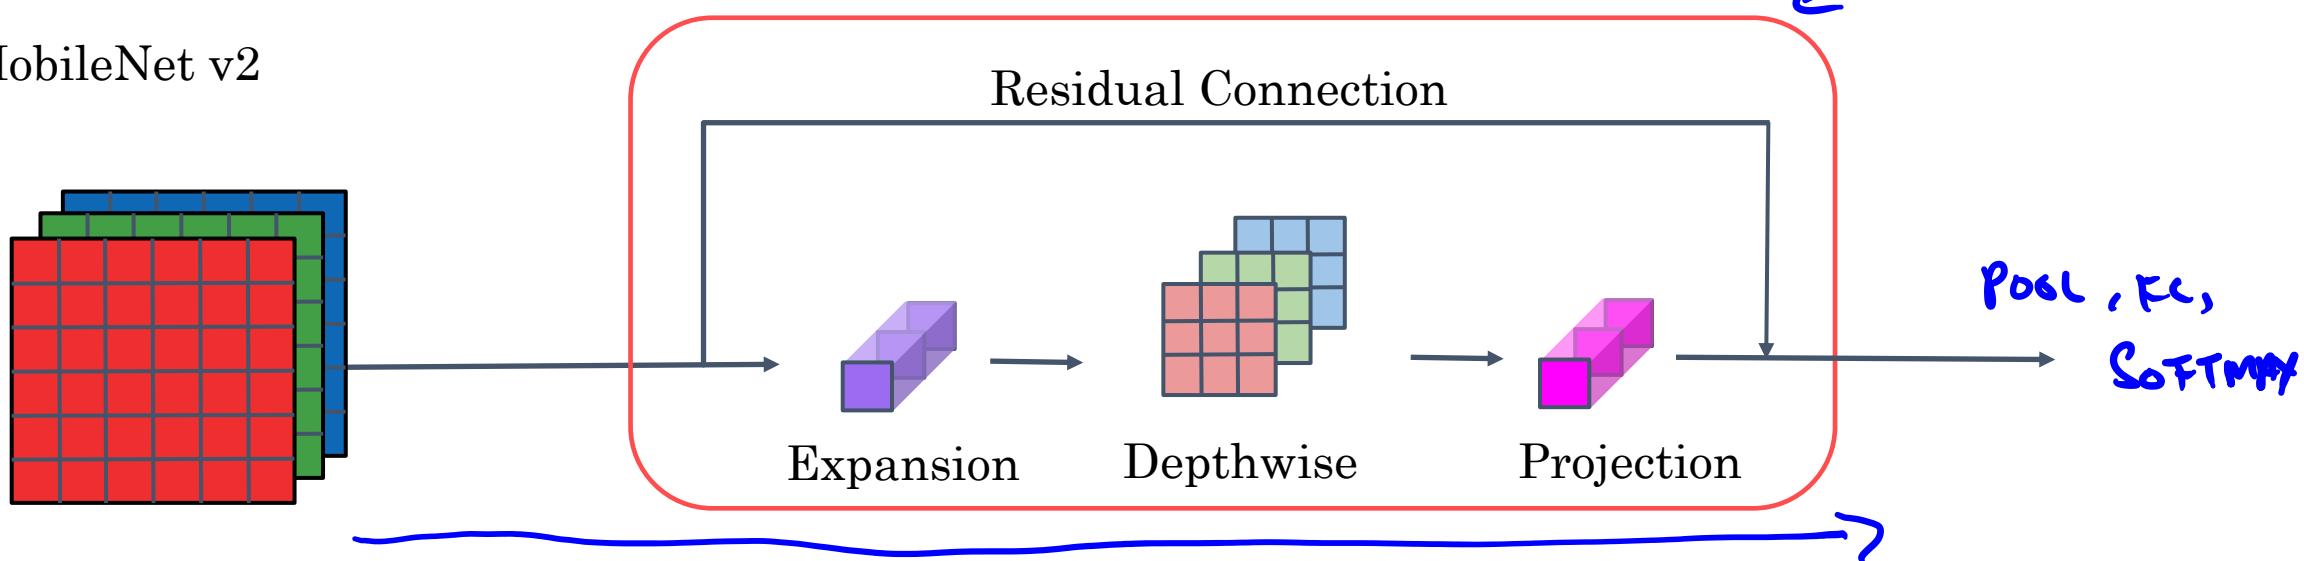


# MobileNet

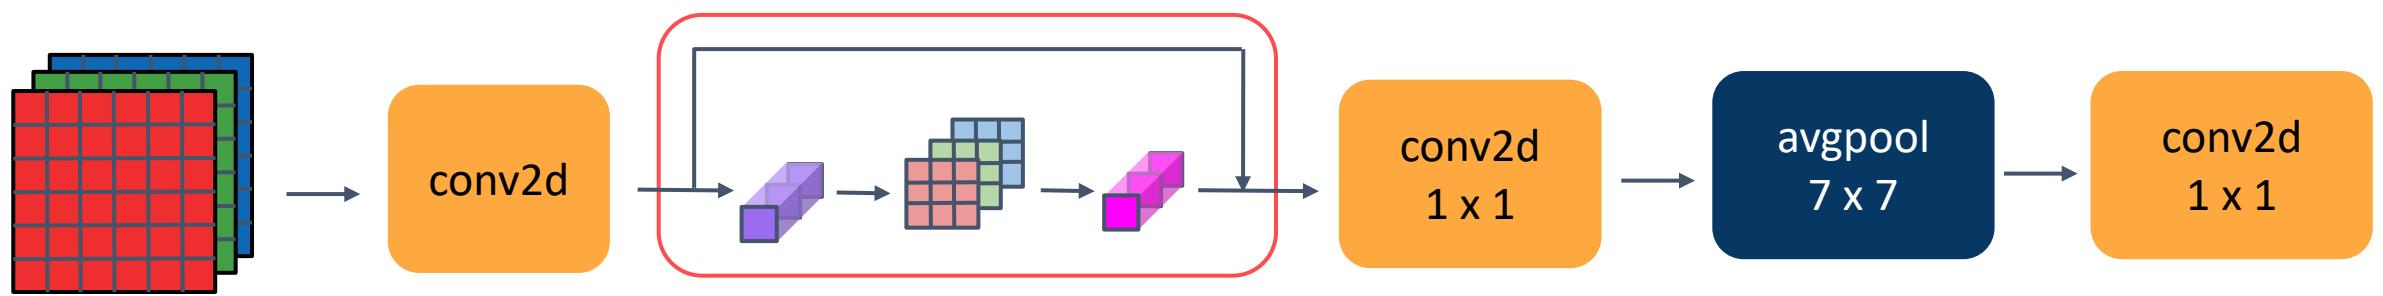
MobileNet v1



MobileNet v2



# MobileNet v2 Full Architecture





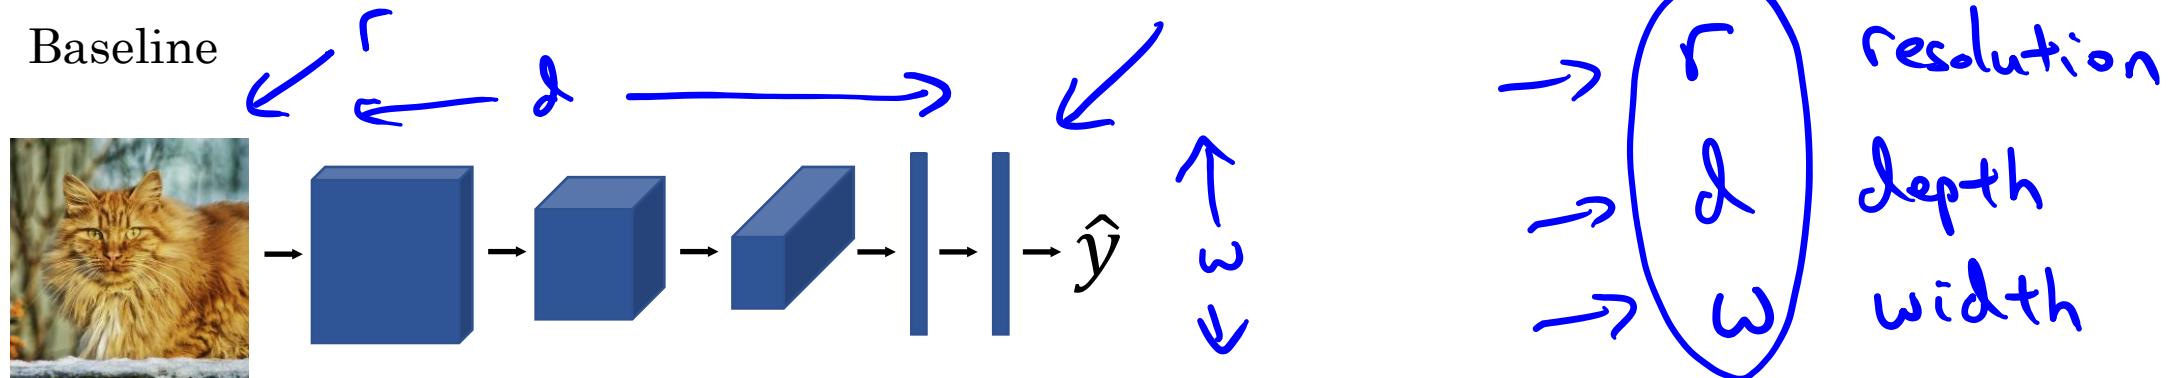
deeplearning.ai

# Convolutional Neural Networks

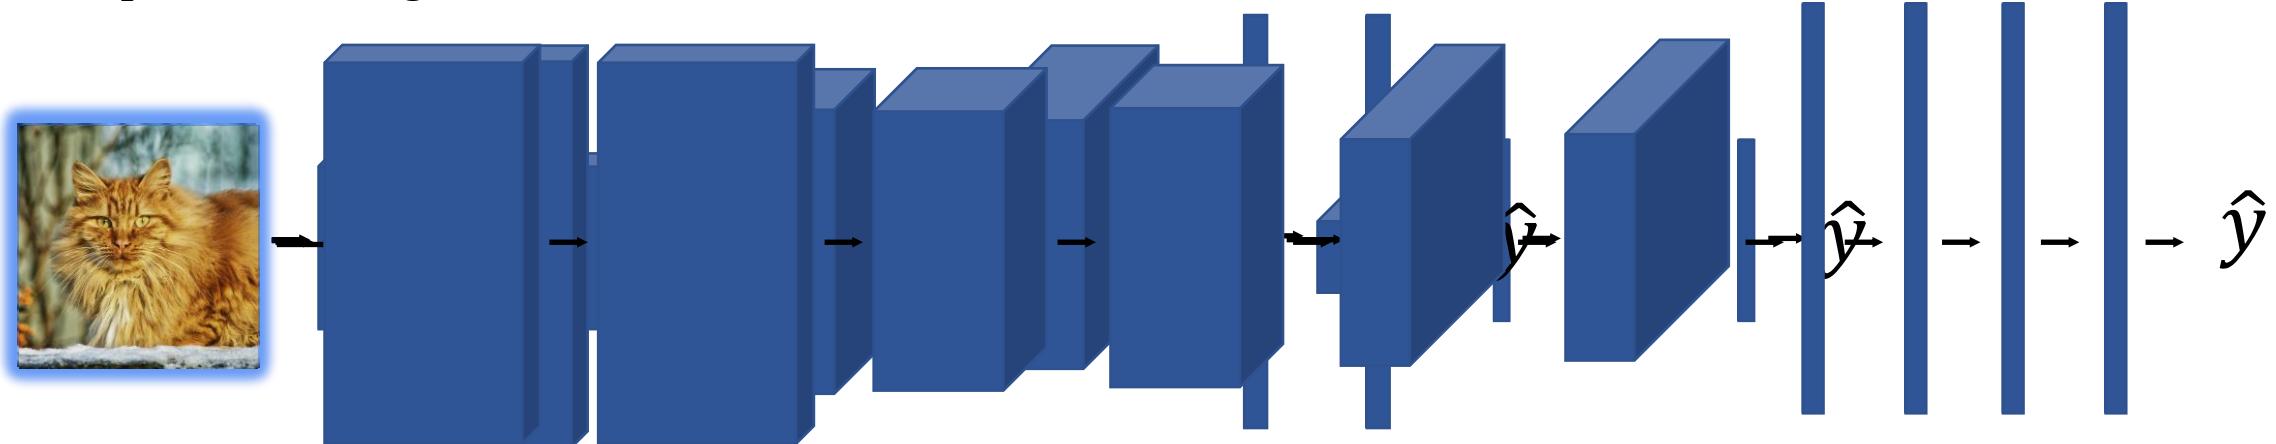
---

## EfficientNet

# EfficientNet



Widthwise Rescaling





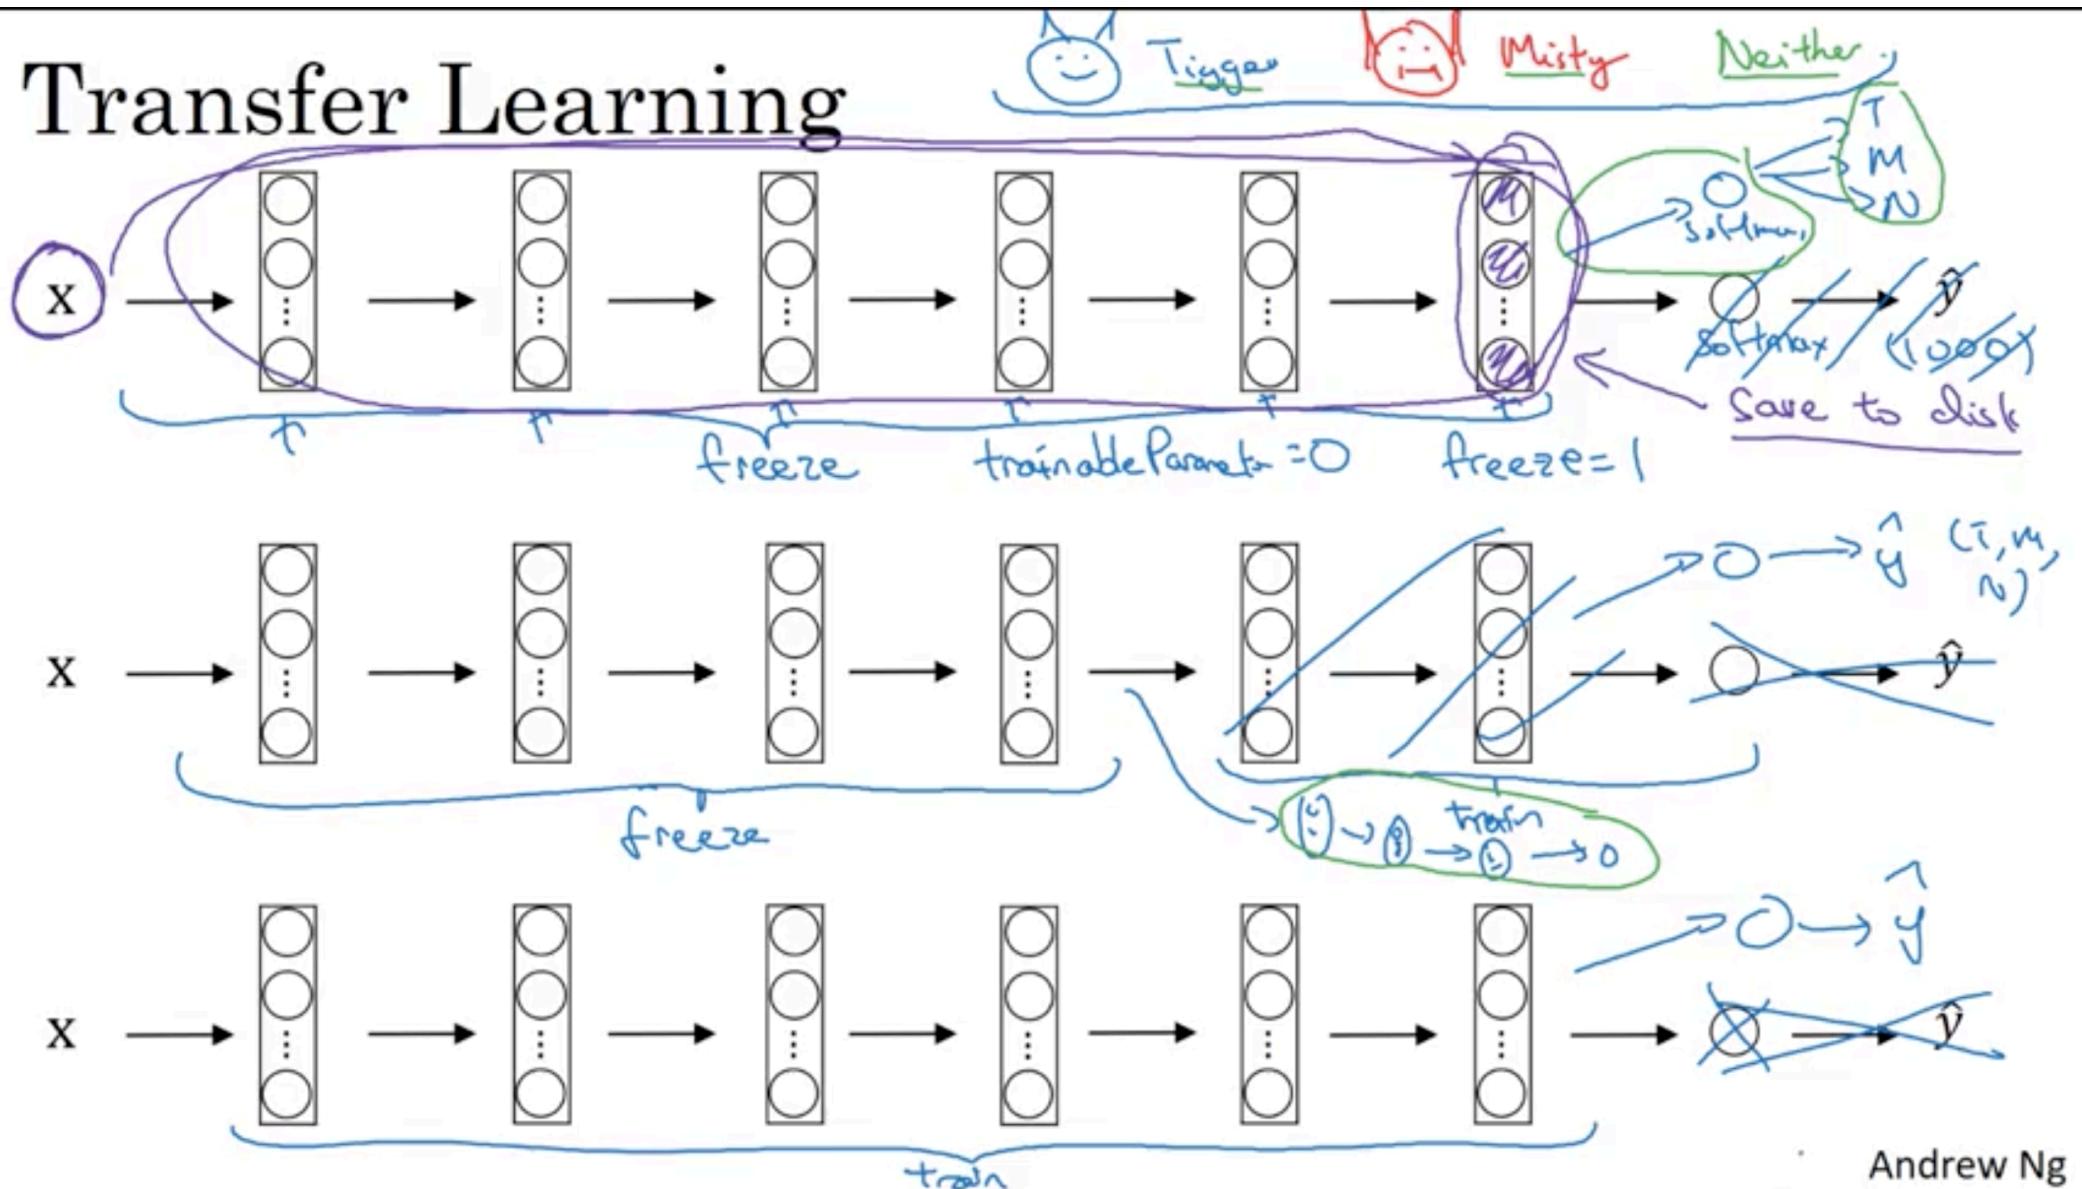
deeplearning.ai

Practical advice for  
using ConvNets

---

Transfer Learning

# Transfer Learning





deeplearning.ai

Practical advice for  
using ConvNets

---

Data augmentation

# Common augmentation method

Mirroring



yc

Random Cropping

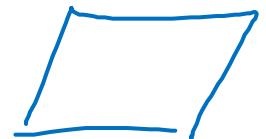


Rotation

Shearing

Local warping

...



# Color shifting



y

R G B  
↓ ↓ ↓  
+20,-20,+20



-20,+20,+20

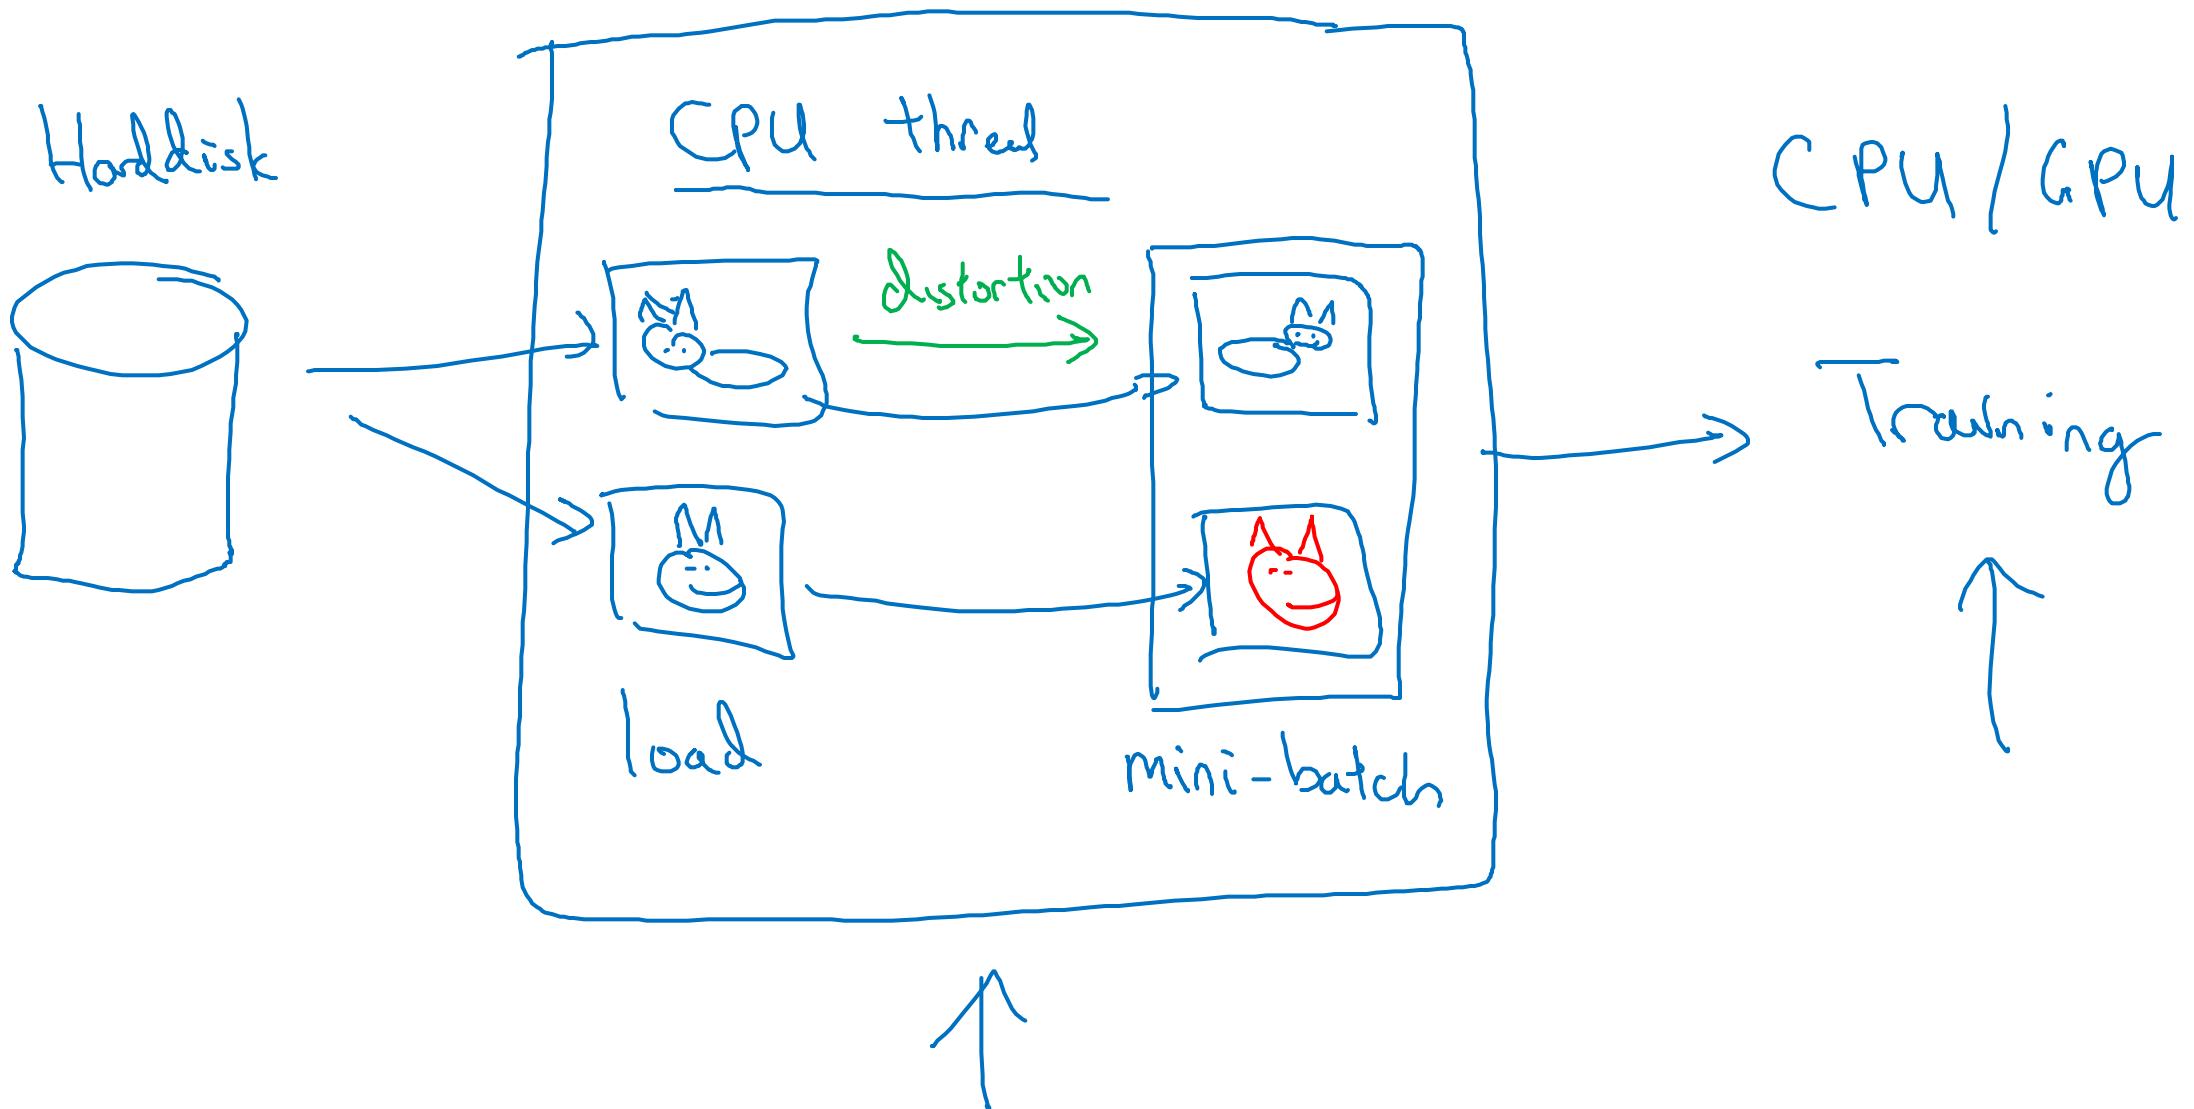


+5,0,+50



Advanced:  
PCA  
[ml-class.org](http://ml-class.org)  
[ AlexNet paper  
[ "PCA color augmentation."  
R B      G

# Implementing distortions during training





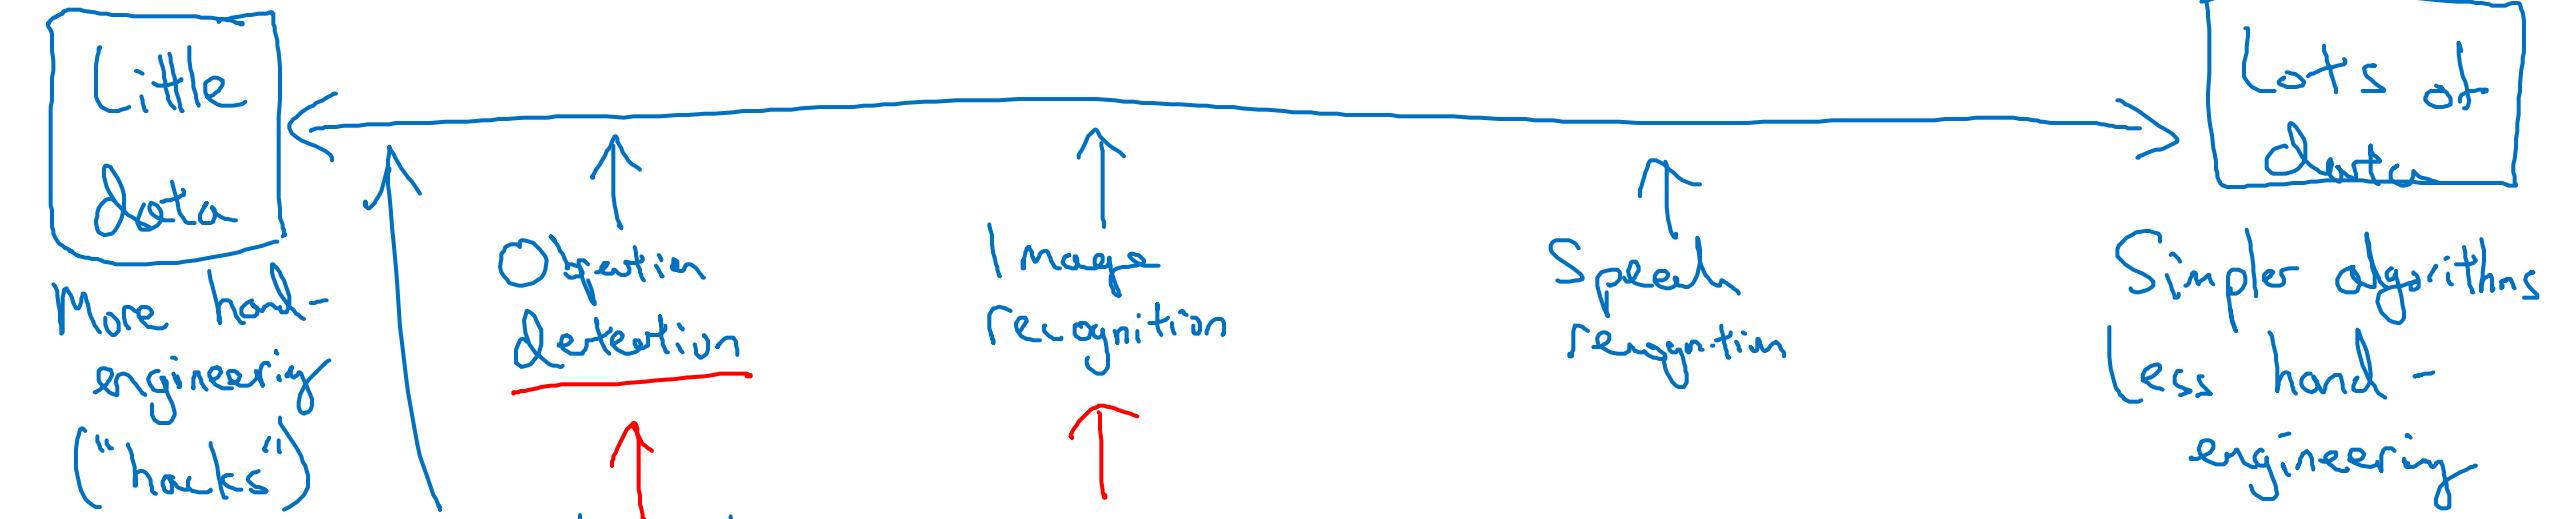
deeplearning.ai

Practical advice for  
using ConvNets

---

The state of  
computer vision

# Data vs. hand-engineering



Two sources of knowledge

- • Labeled data  $(x, y)$
- • Hand engineered features/network architecture/other components



# Tips for doing well on benchmarks/winning competitions

## Ensembling

3 - 15 networks

$\rightarrow \hat{y}$

- Train several networks independently and average their outputs

## Multi-crop at test time

- Run classifier on multiple versions of test images and average results

10-crop



1



+

4



1



+

4

# Use open source code

- Use architectures of networks published in the literature
- Use open source implementations if possible
- Use pretrained models and fine-tune on your dataset